



THE AGENTIC AI REVOLUTION IN SUSTAINABLE TOURISM: A RETRIEVAL-
AUGMENTED GENERATION (RAG) FRAMEWORK FOR COMMUNITY-BASED
KNOWLEDGE HUBS IN THAILAND

MR. SOMATAT NA TAKUATUNG

A THESIS SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE DEGREE OF DOCTOR OF PHILOSOPHY (DESIGN AND PLANNING)
SCHOOL OF ARCHITECTURE AND DESIGN
KING MONGKUT'S UNIVERSITY OF TECHNOLOGY THONBURI

2025

The agentic ai revolution in sustainable tourism: a Retrieval-Augmented Generation
(RAG) framework for community-based knowledge hubs in Thailand

Mr. Somatat Na Takuatung M. Design (Human Centered Design)

A Thesis Submitted in Partial Fulfillment of the Requirements
for the Degree of Doctor of Philosophy (Design and Planning)
School of Architecture and Design
King Mongkut's University of Technology Thonburi
2025

Thesis Committee

..... (Prof. Dr. Pattanasak Mongkolwat, Ph.D.)	Chairman of Thesis Committee
..... (Assoc. Prof. Dr. Chokeanand Bussracumpakorn, Ph.D.)	Member and Thesis Advisor
..... (Asst. Prof. Dr. Chanyaporn Bstieler, Ph.D.)	Member
..... (Assoc. Prof. Dr. Paima Israsena Na Ayudhya, Ph.D.)	Member
..... (Assoc. Prof. Dr. Paniti Netinant, Ph.D.)	Member

Thesis Title	The agentic ai revolution in sustainable tourism: a Retrieval-Augmented Generation (RAG) framework for community-based knowledge hubs in Thailand
Thesis Credits	36
Candidate	Mr. Somatat Na Takuatung
Thesis Advisor	Assoc. Prof. Dr.Chokeanand Bussracumpakorn
Program	Doctor of Philosophy
Field of Study	Design and Planning
Faculty	School of Architecture and Design
Academic Year	2025

ABSTRACT

Thailand's tourism sector is a vital economic driver, yet small-to-medium boutique hotels often struggle to provide personalized, dynamic local recommendations due to limited human resources and technological constraints. Traditional recommendation systems and generic Large Language Models (LLMs) frequently fail in these environments, suffering from "cold-start" problems and "hallucinations"—the fabrication of services that do not exist. Concurrently, relying solely on human agents limits scalability and real-time responsiveness when managing volatile, unstructured community data.

This research addresses these limitations by proposing a "Hotel, Community, and Tourist Relationship Platform" driven by Agentic Artificial Intelligence and Retrieval-Augmented Generation (RAG). The primary objective is to position boutique hotels as central digital knowledge hubs that foster a symbiotic relationship between tourists and local community services. Utilizing a constructive research methodology rooted in the real-world context of Koh Lanta, Krabi, the study digitizes culturally rich community data into a searchable vector database using the Firebase Genkit framework.

To validate the system, the research conducted a dual-comparative evaluation. A synthetic dataset of 50 unique user personas was generated to simulate 150 diverse tourist queries. The primary contribution of this study is a new architectural framework that validates the use of Agentic RAG in unstructured micro-economies, coupled with a new theoretical taxonomy for identifying AI hallucinations in hospitality. The Agentic RAG system was benchmarked against both a standard Base LLM and human domain experts using the Ragas framework to assess Faithfulness, Relevancy, and Safety.

The analysis shows that the proposed RAG Agent achieved a 124% increase in Faithfulness (0.56 vs. 0.25) compared to the base LLM, completely eliminating unsafe recommendations (0.00 maliciousness score). Furthermore, when compared to human agents, the RAG system demonstrated superior consistency, faster data retrieval across complex constraints, and an equal capacity for aligning recommendations with

sustainable tourism principles. This thesis concludes by defining the "Silicon Concierge," a new professional paradigm where AI serves as a trustworthy, scalable bridge between global travelers and local Thai culture, augmenting human capabilities rather than merely replacing them.

Keywords : Agentic AI / Retrieval-Augmented Generation (RAG) / Human Agents / Sustainable Tourism / Community-Based Tourism / Boutique Hotels / Hallucination Mitigation

หัวข้อวิทยานิพนธ์	การปฏิวัติเอไอเชิงตัวแทนในการท่องเที่ยวอย่างยั่งยืน: กรอบการทำงานการสร้างแบบเสริมการสืบค้น (RAG) สำหรับศูนย์กลางความรู้ชุมชนในประเทศไทย
หน่วยกิต	36
ผู้เขียน	นายโสมทัต ฌ ตะกั่วทุ่ง
อาจารย์ที่ปรึกษา	รศ. โชคอนันต์ บุษราคัมภากร
หลักสูตร	ปรัชญาดุษฎีบัณฑิต
สาขาวิชา	การออกแบบและวางแผน
คณะ	สถาปัตยกรรมศาสตร์และการออกแบบ
ปีการศึกษา	2568

บทคัดย่อ

ภาคการท่องเที่ยวของประเทศไทยถือเป็นกลไกสำคัญในการขับเคลื่อนเศรษฐกิจ อย่างไรก็ตาม โรงแรมบูติกขนาดกลางและขนาดย่อมมักประสบปัญหาในการให้คำแนะนำเชิงลึกระดับท้องถิ่นแบบเรียลไทม์ เนื่องจากข้อจำกัดด้านทรัพยากรบุคคลและเทคโนโลยี โมเดลภาษาขนาดใหญ่ (LLMs) ทั่วไปมักล้มเหลวในสภาพแวดล้อมเช่นนี้ โดยประสบปัญหา "อาการประสาทหลอนของเอไอ" (Hallucinations) ในขณะเดียวกัน การพึ่งพาเพียงพนักงานที่เป็นมนุษย์ (Human Agents) ก็มีข้อจำกัดด้านความเร็วและการจัดการข้อมูลชุมชนที่มีความผันผวนสูง

งานวิจัยฉบับนี้จึงนำเสนอ “แพลตฟอร์มความสัมพันธ์ระหว่างโรงแรม ชุมชนและนักท่องเที่ยว” รูปแบบใหม่ที่ขับเคลื่อนด้วยเอไอเชิงตัวแทน (Agentic AI) และเทคโนโลยีการสร้างแบบเสริมการสืบค้น (RAG) โดยมีวัตถุประสงค์เพื่อยกระดับให้โรงแรมบูติกกลายเป็นศูนย์กลางความรู้ดิจิทัล งานวิจัยนี้ใช้ระเบียบวิธีวิจัยเชิงสร้างสรรค์ อ้างอิงบริบทจริงจากพื้นที่เกาะลันตา จังหวัดกระบี่ ในการแปลงข้อมูลชุมชนให้เป็นฐานข้อมูลเวกเตอร์ผ่านเฟรมเวิร์ก Firebase Genkit

เพื่อประเมินประสิทธิภาพของระบบ ผู้วิจัยได้ทำการทดสอบเปรียบเทียบแบบคู่ขนาน (Dual-comparative evaluation) โดยจำลองตัวตนผู้ใช้ 50 รูปแบบ และคำถาม 150 ข้อ คุณูปการหลักของงานวิจัยนี้คือ การนำเสนอกรอบสถาปัตยกรรมใหม่ที่ยืนยันประสิทธิภาพการใช้ระบบ Agentic RAG ในเศรษฐกิจจุลภาคระดับชุมชนที่ไม่มีโครงสร้างตายตัว (Unstructured micro-economies) ควบคู่ไปกับการสร้างอนุกรมวิธานเชิงทฤษฎีแบบใหม่ (New theoretical taxonomy) เพื่อระบุและจำแนกปัญหาข้อมูลบิดเบือนของเอไอ (AI hallucinations) ในอุตสาหกรรมบริการ ทั้งนี้ ระบบ RAG เชิงตัวแทนถูกนำไปเปรียบเทียบกับโมเดลภาษาขนาดใหญ่มาตรฐาน (Base LLM) และเปรียบเทียบกับผู้เชี่ยวชาญที่เป็นมนุษย์ (พนักงานโรงแรม) ผ่านกรอบการวัดผล Ragas เพื่อประเมินความถูกต้องแม่นยำ ความเกี่ยวข้อง และความปลอดภัย

ผลการประเมินแสดงให้เห็นถึงความก้าวหน้าสำคัญ ตัวแทนระบบ RAG มีค่าความถูกต้องแม่นยำ (Faithfulness) เพิ่มขึ้นถึงร้อยละ 124 เมื่อเทียบกับโมเดลทั่วไป และกำจัดการให้คำแนะนำที่ไม่

ปลอดภัยได้อย่างสมบูรณ์ (คะแนนความเป็นอันตราย 0.00) นอกจากนี้ เมื่อเปรียบเทียบกับผู้เชี่ยวชาญที่เป็นมนุษย์ ระบบ RAG ยังแสดงให้เห็นถึงความสม่ำเสมอที่เหนือกว่า ความรวดเร็วในการสืบค้นข้อมูลภายใต้เงื่อนไขที่ซับซ้อน และมีศักยภาพทัดเทียมมนุษย์ในการให้คำแนะนำที่สอดคล้องกับหลักการท่องเที่ยวอย่างยั่งยืน คุณลักษณะฉบับนี้สรุปด้วยแนวคิด "พนักงานต้อนรับดิจิทัล" (Silicon Concierge) ซึ่งเป็นกระบวนการที่เอไอทำหน้าที่เป็นสะพานเชื่อมที่น่าเชื่อถือระหว่างนักท่องเที่ยวและวัฒนธรรมท้องถิ่น โดยมุ่งเน้นการเสริมศักยภาพของมนุษย์มากกว่าการเข้ามาแทนที่

คำสำคัญ : เอไอเชิงตัวแทน / การสร้างแบบเสริมการสืบค้น (RAG) / ผู้เชี่ยวชาญที่เป็นมนุษย์ / การท่องเที่ยวอย่างยั่งยืน / การท่องเที่ยวโดยชุมชน / โรงแรมบูติก / การลดปัญหาข้อมูลบิดเบือนของเอไอ

ACKNOWLEDGEMENT

The completion of this doctoral dissertation would not have been possible without the support, guidance, and encouragement of many individuals and organizations.

First and foremost, I would like to express my deepest gratitude to my principal advisor, Assoc. Prof. Dr. Chokeanand Bussracumpakorn, Ph.D. His visionary guidance, rigorous academic standards, and unwavering support throughout this research journey have been invaluable. His mentorship not only shaped the technical and theoretical direction of this thesis but also profoundly influenced my overall development as a researcher in the field of Design and Planning.

I also wish to extend my sincere appreciation to the faculty members, committee, and staff at King Mongkut's University of Technology Thonburi (KMUTT) for providing an intellectually stimulating environment and the foundational resources required to conduct this extensive research.

A special note of profound gratitude goes to the local community members, boutique hotel owners, and communities in Koh Lanta, Krabi. Their willingness to share their time, operational insights, and deep cultural knowledge formed the very core of the community data utilized in this study. Without their open cooperation and dedication to sustainable practices, the development of the Agentic RAG system would not have been grounded in the authentic realities of local tourism.

Finally, I am eternally grateful to my family and friends for their endless patience, understanding, and moral support during the long and challenging process of completing this Ph.D. Their unwavering belief in my academic endeavors has been a constant source of strength.

TABLE OF CONTENT

ABSTRACT	ERROR! BOOKMARK NOT DEFINED.
LIST OF TABLES	XII
LIST OF FIGURES	XIII
LIST OF ABBREVIATIONS AND TERMS.....	XIV
CHAPTER 1 INTRODUCTION	1
1.1 Background of the Study	1
1.2 Statement of the Problem	2
1.3 Research Objectives	4
1.4 Conceptual Framework of the Study: The Hotel as a Knowledge Hub.....	4
1.5 The Proposed Solution: The Tripartite Knowledge-Sharing Ecosystem.....	6
1.5.1 Conceptualizing the Tripartite Ecosystem	7
1.5.2 Platform Development and Architecture.....	7
1.6 Scope and Delimitations of the Study	10
1.7 Significance of the Study	11
1.8 The Socio-Cultural Context of the Study Area (Koh Lanta, Krabi).....	11
1.9 Summary of Chapter 1	12
CHAPTER 2 LITERATURE REVIEW	13
2.1 The Evolution of Digitalization in Hospitality	13
2.1.1 Pre-Pandemic Technology and the "Human Touch"	14
2.1.2 COVID-19 and the Acceleration of Contactless Systems.....	14
2.1.3 The Post-Pandemic Macro-Economic Imperative	15
2.1.4 Key driving factors for AI.....	15
2.2 Data Analytics in the Service Industry.....	17
2.2.1 What is Data Analytics?.....	17
2.2.2 Hotel Data Types.....	18
2.3 Hotel Data Analytics	20
2.3.1 Customer Segmentation	21

2.3.2	Customer Profiling	21
2.3.3	Menu Engineering	21
2.3.4	Productivity Indexing	22
2.3.5	Customer Associations and Sequencing	22
2.3.6	Forecasting	23
2.3.7	Energy Consumption.....	23
2.3.8	The Right Room at the Right Rate	23
2.4	Traditional Recommendation Systems and Their Limitations.....	24
2.4.1	Principles of Community-Based Tourism (CBT)	24
2.4.2	The Evolution from CSR to CSV	25
2.4.3	The Shared Value Model in Platform Economies.....	25
2.4.4	The "Cold-Start" Problem in Informal Micro-Economies	26
2.4.5	Typology of Recommendation Systems in Tourism.....	27
2.4.6	Limitations of Traditional Recommender Systems in CBT.....	28
2.5	Large Language Models (LLMs) and the Hallucination Dilemma	28
2.5.1	The Rise of Generative AI in Travel.....	29
2.5.2	Parametric Memory and Domain-Specific Hallucinations	29
2.5.3	Hallucinations in LLM-Based Systems	30
2.5.4	Trustworthy AI.....	30
2.6	Community Data Governance in AI Ecosystems	31
2.7	Retrieval-Augmented Generation (RAG) and Agentic Systems.....	32
2.7.1	The Agentic Paradigm.....	32
2.7.2	Retrieval-Augmented Generation (RAG)	32
2.7.3	The Mechanics of RAG	33
2.7.4	AI Agents and Autonomous Decision Systems	34
2.8	Identification of the Research Gap	34
2.9	Firestore Genkit as an AI Framework.....	35
2.9.1	Recommendation Evaluation Metrics	36
2.10	The "Hotel, Community, and Tourist Relationship Platform"	39
2.11	Summary of Chapter 2	40
CHAPTER 3 METHODOLOGY AND SYSTEM ARCHITECTURE		42
3.1	Research Design and Paradigm	42
3.1.1	Data Governance and PDPA Compliance Framework	43
3.2	System Architecture: The Tripartite Ecosystem.....	44
3.2.1	The Context Retriever and Real-Time Data Ingestion.....	46
3.2.2	The RecommendDotPrompt and Generation Phase.....	48
3.3	Agentic RAG Implementation via Firestore Genkit.....	48
3.3.1	Database Schema and Prompt Engineering	49

3.4 System Evaluation Procedures.....	50
3.4.1 Experimental Setup	50
3.4.2 Data Synthesis and Persona Generation.....	50
3.4.3 Justification for Synthetic Data Utilization.....	50
3.4.4 Dataset Preparation and Chunking.....	52
3.4.5 Persona Construction and Diversity	54
3.5 Automated Evaluation via Firebase Genkit and Ragas.....	55
3.6 Manual Validation via Human Domain Experts.....	57
3.6.1 Panel Selection and Expertise	58
3.6.2 Blind Review Procedure	59
3.6.3 Evaluation Criteria	59
3.7 Summary of Chapter 3	60
CHAPTER 4 RESULTS AND EVALUATION	62
4.1 Empirical Grounding via Unsupervised K-Means Clustering.....	62
4.2 Taxonomy of Tourism Hallucinations.....	63
4.3 Selection Criteria for Qualitative Case Studies.....	64
4.4 Qualitative Analysis: Side-by-Side Case Studies.....	64
4.4.1 Case Study 1 The Cold-Start and Short-Stay Dilemma	65
4.4.2 Case Study 2 Eradicating Fabricated Entities in Niche CBT.....	65
4.4.3 Case Study 3 Mitigating Attribute Misalignment in Capacity Constraints.....	66
4.4.4 Case Study 4 Culturally Specific Semantic Fulfillment	67
4.4.5 Case Study 5 Eco-Tourism and Sustainable Routing	67
4.4.6 Case Study 6 Overcoming Temporal Hallucinations.....	68
4.4.7 Case Study 7 Eliminating Unverified Transportation (Safety)	68
4.5 Quantitative Results Ragas Framework Benchmarking.....	69
4.5.1 Faithfulness (Factual Grounding)	70
4.5.2 Answer Relevancy (Semantic Fulfillment).....	70
4.5.3 Maliciousness (Safety Validation)	71
4.5.4 Analysis of System Safety and Maliciousness Mitigation.....	72
4.6 Validation through Human Expert Evaluation.....	72
4.6.1 Evaluation Methodology	73
4.6.2 Comparative Results and Expert Findings	75
4.7 Validation of Qualitative Heuristics via Inter-Rater Reliability	76
4.7.1 Conclusion of the Evaluation	76
4.8 Interpretation of Results and System Viability	77
4.9 Summary of Chapter 4	77
CHAPTER 5 CONCLUSION AND RECOMMENDATIONS.....	79

5.1 Conclusion and Summary of Contributions	79
5.2 Practical and Managerial Implications	81
5.3 Real-World Deployment Governance	82
5.4 Limitations of the Study	83
5.5 Future Work and Scalability.....	84
5.6 Concluding Remarks	89
5.7 Summary of Chapter 5	89
REFERENCES	90
APPENDIX	16
Appendix A system data structures and knowledge corpora.....	17
A.1 Overview of the Knowledge Base.....	17
A.2 Tourist Persona Dataset (Guest Profiles)	17
A.3 Community-Based Tourism (CBT) Dataset.....	17
A.4 Automated Genkit Evaluation Inputs (Test Case Batches).....	18
Appendix B: Ragas Framework & Human Expert Evaluation Metrics.....	20
B.1 Overview of the Ragas Scoring Dataset.....	20
B.2 Dual-Evaluation Query Metrics (Sample).....	20
Appendix C: System Prompt Configurations (Dotprompts)	22
C.1 Overview of Agentic Instructions	22
C.2 Generation Phase: The Recommendation Prompt.....	22
C.3 Retrieval Phase: The Semantic Router Prompt	23
Appendix D: Nosql Database Schema And Data Dictionary	26
D.1 Overview of the Firestore Architecture	26
D.2 Collection: hotels.....	26
D.3 Collection: guests	27
D.4 Collection: partners	27

LIST OF TABLES

TABLE	PAGE
2.1 Hotel Data Types.	19
4.1 K-Means Cluster Centroid (k=3) based on Synthetic Persona Queries	63
4.2 Comparative Output Analysis for Mitigating the Cold-Start Dilemma.	65
4.3 Comparison of System Outputs Mitigating Fabricated Entity Hallucinations.	65
4.4 Comparative Analysis of Physical Safety Constraints and Maliciousness Migration.	66
4.5 Evaluation of Culturally Specific Semantic Fulfillment and Geographic Accuracy.	67
4.6 Comparative Analysis of Eco-Tourism and Sustainable Routing.	67
4.7 Qualitative Benchmarking: Human Expert Evaluation of System Viability.	68
4.8 Accuracy Comparison for Retrieving Dynamic Local Transfer Services.	69
4.9 Sample Evaluation Log Showing Scores for RAG Agent vs. Base LLM.	69
4.10 Comparative Performance Summary of Base LLM vs. Agentic RAG across RAGAS Metrics.	69
4.11 Sample User Queries and Evaluation Scores.	73
4.12 Aggregate Evaluation Results (N=50).	74

LIST OF FIGURES

FIGURE	PAGE
1.1 Conceptual of The Hotel-Community-Tourist Ecosystem, “The Koh Lanta Shared Value Model”.	6
1.2 Hotel Front App: Dashboard for Hotel A.	8
1.3 Community App: Event Calendar and Order Tracking for Local Vendor.	9
1.4 Mobile Guest App: Login, Home, and Recommendations.	9
2.1 Key Drivers for AI.	17
2.2 Hotel Data Map.	18
3.1 Architectural of the Application and the Recommendation Agent.	46

LIST OF ABBREVIATIONS AND TERMS

Agentic AI An advanced artificial intelligence framework where the system functions as an autonomous agent capable of reasoning, planning, and executing multi-step tasks using external tools, moving beyond static text generation (Sivan, 2024).

AI Hallucination A phenomenon where a generative model produces false, fabricated, or nonsensical information presented confidently as fact. This study specifically categorizes tourism hallucinations into *Fabricated Entities*, *Attribute Misalignment*, and *Temporal Decay*.

Base LLM (Base Large Language Model) A neural network-based text generation model that relies exclusively on its pre-trained, static parametric memory.

CSV (Creating Shared Value) A strategic business framework introduced by Porter and Kramer (2011) that focuses on generating economic value in a way that simultaneously produces value for society by addressing its challenges. In this research, CSV is operationalized through the AI architecture, which enhances the guest experience for boutique hotels while programmatically driving economic visibility and tourism revenue to local micro-enterprises.

CBT (Community-Based Tourism) A form of sustainable tourism owned and managed by the local community, aimed at generating economic benefits for residents while preserving cultural heritage (Moscardo, 2023).

Cold-Start Problem A limitation in traditional recommendation systems where the engine cannot draw accurate inferences for new users or new items due to an absence of historical interaction data (Yuan & Hernandez, 2023).

Dotprompt A feature within the Firebase Genkit framework that allows prompts to be written, managed, and executed as distinct code files, ensuring consistent generation formats.

Genkit An open-source framework developed by Firebase designed to simplify the integration of AI-powered features and orchestrate the interaction between Google Cloud services, vector databases, and Large Language Models.

LLM (Large Language Model) A deep learning algorithm that can recognize, summarize, predict, and generate text based on knowledge gained from massive pre-training datasets.

RAG (Retrieval-Augmented Generation) An AI architecture that enhances generative models by fetching facts from an external, authoritative knowledge base before generating a response, thereby mitigating hallucinations (Lewis et al., 2020).

Ragas Framework A standardized evaluation framework used to objectively quantify and measure the performance of RAG applications across specific metrics, such as Faithfulness, Answer Relevancy, and Maliciousness.

SAR (Sustainability-Augmented Reranking) A theoretical mechanism proposed for future system development that explicitly weights AI recommendations based on environmental metrics, such as carbon footprints, to prioritize sustainable travel options.

CHAPTER 1 INTRODUCTION

1.1 Background of the Study

The tourism and hospitality sector serves as a vital cornerstone of Thailand's gross domestic product (GDP) and a primary driver of national economic growth. Known globally as a premier travel destination, Thailand benefits from a diverse array of attractions ranging from the metropolitan hub of Bangkok to the ecotourism landscapes of the North and the coastal communities of the South.

However, the unprecedented disruption caused by the COVID-19 pandemic served as a critical inflection point for the industry. The ensuing crisis generated severe negative economic impacts across the global and local tourism sectors, forcing small-to-medium enterprises, particularly boutique hotels, to rapidly adapt to a "new normal" characterized by reduced workforce capacities and heightened traveler expectations regarding health and safety.

Initially, this adaptation manifested as reactive digital transformation. To mitigate face-to-face interactions, hotels implemented "contactless" software solutions, enabling self-service check-ins, digital room keys, and mobile ordering. While these tools successfully addressed immediate operational constraints and built upon earlier progress in hospitality information and communication technologies (Law, Buhalis, & Cobanoglu, 2014), they inadvertently diminished the personalized "human touch" that traditionally defines luxury and Thai hospitality. As Subramony et al. (2017) extensively argue, leveraging the human touch in service interactions is critical; over-automation risks commoditizing the guest experience and lowering overall satisfaction.

As the industry moves into the post-pandemic era, the macro-economic directive has shifted. According to Chutijirawong and Sangmanacharoen (2025) note in the Deloitte Thailand Economic Outlook, sustainable recovery relies on modernizing the sector to attract high-value tourism rather than sheer volume. It is no longer sufficient to merely digitize basic transactions. Hotels must now leverage advanced technologies such as Machine Learning (ML) and Artificial Intelligence (AI) to actively manage customer profiling, forecasting, and personalized engagement (Quanovo Data Technologies, 2017). The transition from reactive contactless software to proactive Generative AI systems represents the next frontier in hospitality management, promising to restore the personalized guest experience without proportionally increasing human labor costs (Robinson, 2024).

1.2 Statement of the Problem

While Generative AI and Large Language Models (LLMs) offer unprecedented opportunities for hyper-personalized guest engagement, their deployment in niche, highly localized environments faces significant technical and structural barriers. This is particularly evident in Community-Based Tourism (CBT) destinations like Koh Lanta, Krabi, where community involvement and local knowledge are paramount to sustainable tourism development (Moscardo, 2008).

While the socio-economic goals of CBT align closely with the principles of Creating Shared Value (CSV), the integration of social impact directly into core business operations (Porter & Kramer, 2011), a significant research gap remains in operationalizing CSV within digital, platform-based economies. Current CSV literature lacks frameworks addressing "algorithmic inequality." Traditional digital tourism platforms utilize popularity-based search algorithms that inherently favor massive,

established corporations. This systemic bias marginalizes local micro-enterprises and prevents the equitable distribution of digital traffic to the community.

Standard recommendation systems rely heavily on static databases and historical user patterns, whereas CBT environments are fundamentally dynamic and unstructured. The "items" being recommended, such as a specific longtail boat captain's availability, seasonal local markets, or informal cultural events, are highly volatile and rarely documented in standard digital booking engines. This data volatility, combined with algorithmic inequality, introduces two major failure points for existing technologies.

1. The Cold-Start Problem in Traditional Recommenders

Static models powered by standard collaborative filtering and neural retrieval cannot effectively match travelers to local experiences when the services themselves are ephemeral and lack sufficient historical rating data (Wei et al., 2017). A traveler staying for a short 2-day period requires immediate, highly accurate recommendations, which these traditional platforms cannot generate without prior historical overlap or highly optimized digital profiles. As a result, the algorithm traps the tourist in generic recommendations, actively preventing the creation of shared value.

2. The Hallucination of General-Purpose LLMs

The advent of the Transformer architecture (Vaswani et al., 2017) led to powerful LLMs. However, when standard Base LLMs are queried about niche local services, they rely solely on their pre-trained parametric memory. Lacking access to real-time, hyperlocal data, these generic models frequently "hallucinate," confidently fabricating venues, quoting outdated prices, or suggesting unsafe activities. The critical issues of model faithfulness (Maynez et al., 2020) and toxicity or harm (Rawte et al., 2023) render ungrounded LLMs unsuitable for standalone deployment. In the context of travel, where

factual accuracy dictates user safety, these unmitigated hallucinations violate international guidelines for trustworthy AI (European Commission, 2021; OECD, 2022).

Consequently, a critical gap exists between the advanced capabilities of Generative AI and the unstructured realities of Community-Based Tourism. There is an urgent need for a specialized architectural solution that can retrieve verifiable, real-time community data before generating recommendations.

1.3 Research Objectives

This research aims to construct and validate an Agentic RAG-based recommendation system within a real-world tourism setting. The specific objectives are:

- To design and implement a tripartite "Hotel, Community, and Tourist Relationship Platform" utilizing the Firebase Genkit framework (Firebase, n.d.) to manage complex RAG processes.
- To evaluate the performance of the Agentic RAG system against traditional Base Large Language Models, measuring its ability to eliminate hallucinations and adapt to dynamic local data.
- To measure output quality across standardized AI evaluation metrics (Ragas Team, n.d.), specifically: Faithfulness (factual grounding), Answer Relevancy (semantic fulfillment), and Maliciousness (safety validation).
- To demonstrate the economic and sustainable viability of Agentic AI in creating a mutually beneficial ecosystem for boutique hotels and local communities in Thailand.

1.4 Conceptual Framework of the Study: The Hotel as a Knowledge Hub

The conceptual foundation of this research represents a paradigm shift in how small-to-medium boutique hotels operate within a local ecosystem. Traditionally, hotels utilize

digital tools strictly for internal operational efficiency (e.g., contactless check-ins), inadvertently isolating the guest from the surrounding community.

This study introduces a novel conceptual model that repositions the boutique hotel as a dynamic "Knowledge Hub". In this model, the hotel ceases to be merely a provider of accommodation and instead acts as a centralized socio-technical conduit.

By leveraging Agentic Generative AI, the hotel bridges the gap between unstructured tourist demand (complex, highly personalized traveler preferences) and decentralized community supply (local communities, longtail boat captains, and artisans). Ultimately, this conceptual model aims to facilitate a "Shared Business Value Model," ensuring that technological modernization within the hotel directly decentralizes economic benefits to the broader Koh Lanta community.

Ultimately, this conceptual framework introduces the **Koh Lanta Shared Value Model** (Figure 1.1), a localized socio-economic framework that operationalizes the principles of Creating Shared Value (CSV) (Porter & Kramer, 2011) through a decentralized, community-driven artificial intelligence architecture. Unlike traditional, extractive tourism platforms that rely on biased, review-centric algorithms (Van Dijck et al., 2018), this model utilizes an Agentic Retrieval-Augmented Generation (RAG) system (Lewis et al., 2020; Wang et al., 2024) to ensure equitable market access for local micro-enterprises.

The framework functions upon a reciprocal data ecosystem: it crowdsources live, hyper-local safety and operational constraints directly from community stakeholders, specifically hotel administrators and local vendors, to ground the AI's contextual understanding. In return, the system employs semantic vector matching to bypass the

algorithmic 'cold-start' problem (Wei et al., 2017), connecting tourists with mathematically validated, highly relevant, and safe local experiences. Ultimately, the Koh Lanta Shared Value Model demonstrates that integrating strict, community-governed data constraints into a generative AI pipeline intrinsically aligns computational safety with sustainable, localized economic growth (Moscardo, 2008).




Figure 1.1: Conceptual Model of The Hotel-Community-Tourist Ecosystem, “The Koh Lanta Shared Value Model”. *Source: Author (2025)*

1.5 The Proposed Solution: The Tripartite Knowledge-Sharing Ecosystem

To address algorithmic inequality and operationalize the Koh Lanta Shared Value Model, this research develops a localized digital prototype. This system functions as a Tripartite knowledge-sharing ecosystem connecting hotels, local communities, and tourists. Driven by Agentic Artificial Intelligence (AI) and Retrieval-Augmented Generation (RAG), the proposed community-based prototype shifts the paradigm of the boutique hotel from a

mere accommodation provider into a central digital knowledge hub that fosters a symbiotic economic relationship between tourists and the local community.

1.5.1 Conceptualizing the Tripartite Ecosystem

Unlike standard Generative AI, which primarily focuses on passive content creation, Agentic AI acts as an autonomous planner capable of executing complex workflows, tool-use, and goal-oriented decisions (Wang et al., 2024; Xi et al., 2025). To deploy this in a community-based context, the platform establishes a direct economic bridge. The RAG architecture successfully spans the gap between fluid text generation and verifiable factual data (Gao et al., 2023; Lewis et al., 2020), mathematically ensuring tourists receive safe, data-grounded recommendations. Simultaneously, the local community benefits from direct economic engagement, and the hotel optimizes revenue while reducing administrative staffing expenses.

1.5.2 Platform Development and Architecture

To operationalize this ecosystem, the researcher developed a centralized data collection and routing platform built upon a modern Angular frontend and orchestrated via a Firebase Genkit architecture. This framework features three integrated applications:

1. Hotel Front App

A cloud-based web application designed for hotel staff to manage bookings, guests, rooms, and services, complete with integrated communication features to reduce administrative overhead.

2. Community App

A direct digital conduit linking the hotel, community partners, and guests. Local communities, tour operators, and artisans utilize this application to seamlessly sync their dynamic service offerings, such as pricing and physical safety capacities, directly into the hotel's centralized NoSQL database.

3. Mobile Guest App

A direct-to-consumer interface that streamlines the guest experience. It acts as a personalized travel companion, allowing users to input specific travel constraints and receive personalized Agentic AI recommendations that bypass traditional algorithmic popularity biases.




Figure 1.2: Hotel Front App: Dashboard for Hotel A. *Source: Author (2025)*




Figure 1.3: Community App: Event Calendar and Order Tracking for Local Communities. *Source: Author (2025)*




Figure 1.4: Mobile Guest App: Login, Home, and Recommendations. *Source: Author (2025)*

The Recommendation Agent Workflow The core innovation of this platform lies in its recommendation generation workflow. As hotel staff and community partners manage their respective profiles and services, guests register and provide their specific activity interests. When a guest seeks a personalized itinerary, a dedicated recommendation agent (operating as a cloud API service) is triggered.

This agent performs an embedding search within a centralized vector database across five key tables: hotel services, community services, guest activities, historical orders, and service reviews. The system retrieves the top-k contextually similar chunks of data. These retrieved real-time contexts, alongside the user's specific query, are then used to prompt the Large Language Model (LLM).

To ensure strict control over the AI's output, tone, format, and length, the system utilizes Firebase Genkit's Dotprompt framework (Firebase, n.d.). Managing prompts as discrete code allows the AI to systematically process the retrieved localized information, thereby generating highly personalized, specific recommendations that are completely grounded in the community's real-time reality.

1.6 Scope and Delimitations of the Study

This study is geographically bounded to the island of Koh Lanta, Krabi Province, Thailand. The technological scope encompasses the design, development, and pre-deployment evaluation of the Tripartite knowledge-sharing ecosystem. Built upon a modern Angular frontend and orchestrated via a Firebase Genkit architecture, the research focuses specifically on utilizing Agentic Retrieval-Augmented Generation (RAG) to process real-time community-based tourism (CBT) data, thereby addressing the algorithmic "cold-start" problem for short-stay tourists.

A primary delimitation of this study is its boundary regarding live human-subject testing. Deploying an unverified Generative AI recommendation engine to live tourists presents severe liability risks, such as suggesting unsafe maritime activities during monsoon warnings. To strictly comply with international Trustworthy AI safety guidelines regarding physical harm (European Commission, 2021; OECD, 2022) and to adhere to Thailand's Personal Data Protection Act (PDPA) regarding the unauthorized processing of empirical guest profiles (Personal Data Protection Committee, 2019), the evaluation phase was intentionally delimited to a controlled, simulated environment.

The system's performance was rigorously stress-tested using an empirically constructed dataset of 50 synthetic traveler personas (yielding 150 complex test queries) rather than utilizing live, longitudinal data or historical booking logs from actual Koh Lanta hotel guests. Consequently, the evaluation scope is delimited to the assessment of factual faithfulness, semantic relevancy, and safety (maliciousness) using the automated Ragas framework, alongside qualitative corroboration by a panel of local industry experts. Because this study serves as a foundational pre-deployment validation, it does not encompass longitudinal live-user testing, nor does it measure direct financial conversion

rates from the generated itineraries. The empirical business metrics and real-world deployment of the Koh Lanta Shared Value Model are relegated to future Phase 2 research.

1.7 Significance of the Study

This dissertation offers vital contributions to both academic literature and practical industry application:

1. Theoretical Contribution

It advances the field of applied artificial intelligence in hospitality by formally documenting a domain-specific hallucination taxonomy (fabrications, temporal decay, and attribute misalignment). It also provides a validated blueprint for applying Agentic RAG to overcome the "cold-start" problem in recommendation systems (Banerjee et al., 2025).

2. Practical and Managerial Contribution

For the industry, this research provides a scalable operational model for small-to-medium boutique hotels. By acting as digital hubs, these hotels can offer luxury-level concierge services without associated labor costs, directly routing tourism revenue to community-based vendors.

1.8 The Socio-Cultural Context of the Study Area (Koh Lanta, Krabi)

To fully understand the necessity of a localized, Agentic AI platform, one must first understand the complex socio-cultural fabric of the chosen study area. Koh Lanta is an archipelago situated in the Andaman Sea within Thailand's Krabi province. Unlike heavily commercialized hubs such as Phuket or Pattaya, Koh Lanta possesses a unique demographic mosaic that directly shapes its Community-Based Tourism offerings.

Historically, the island's earliest inhabitants were the Urak Lawoi (Sea Gypsies), an indigenous seafaring community with distinct animistic beliefs and a deep connection to the marine ecosystem. Over the centuries, the island evolved into a vital trading port, attracting Thai-Muslim fishermen and Thai-Chinese merchants. Today, approximately 95% of the island's population is of Muslim descent, living harmoniously alongside the remaining ethnic Chinese and Urak Lawoi communities.

Following the Thai government's push for tourism development, Koh Lanta transitioned into an open destination celebrated for its "Green Tourism," focusing on eco-friendly footprints and cultural immersion (e.g., Tung Yee Peng mangrove tours). This rich, multi-layered cultural landscape is precisely why generic algorithms fail here. A standard recommendation engine cannot grasp the cultural significance of local festivals nor dynamically track the seasonal availability of a local family's longtail boat charter. Consequently, this study leverages Koh Lanta as the ultimate stress test for Agentic AI.

1.9 Summary of Chapter 1

This chapter has established the foundational premise of the dissertation: the critical need to bridge the gap between advanced Generative AI and the complex realities of Community-Based Tourism in post-pandemic Thailand. By identifying the limitations of Base LLMs, namely their tendency to hallucinate (Maynez et al., 2020), and the failure of traditional algorithms to overcome the cold-start problem, this study proposes a transformative solution. The "Hotel, Community, and Tourist Relationship Platform," powered by Agentic RAG, aims to elevate the guest experience and foster a sustainable ecosystem for communities like Koh Lanta.

CHAPTER 2 LITERATURE REVIEW

To contextualize the development of the proposed platform, this chapter reviews existing literature across five core themes: (1) the evolution of digitalization in post-pandemic hospitality, (2) the role of data analytics in customer profiling, (3) the mechanisms and limitations of traditional recommendation systems, (4) the emergence of Retrieval-Augmented Generation (RAG) and Agentic AI, and finally, (5) the conceptualization of the Tripartite Ecosystem.

It begins by tracing the evolutionary trajectory of digital transformation within the hospitality sector, heavily influenced by the COVID-19 pandemic. It then examines the socio-economic principles of CBT and the structural barriers that prevent traditional recommendation systems from succeeding in these dynamic environments.

Subsequently, the chapter evaluates the advent of Generative AI, addressing the pervasive issue of algorithmic hallucination. Finally, it explores Agentic RAG as an architectural solution, culminating in the identification of the specific research gap this dissertation addresses.

2.1 The Evolution of Digitalization in Hospitality

Prior to the COVID-19 pandemic, Thailand's tourism sector was a primary engine of national economic growth, accounting for nearly a fifth of the country's gross domestic product (GDP). However, the pandemic exposed the fragility of an industry over-reliant on sheer tourist volume and face-to-face service models. As Chutijirawong and Sangmanacharoen (2025) state in the Deloitte Thailand Economic Outlook, the post-pandemic recovery strategy fundamentally shifts away from mass tourism toward high-value, sustainable, and experiential travel.

For boutique hotels, which lack the massive capital reserves of international chains, this economic pivot presents a dual challenge: they must elevate the personalized "human touch" to attract premium travelers, while simultaneously keeping operational and labor costs strictly controlled. Digital transformation is no longer optional; it is a vital survival mechanism.

2.1.1 Pre-Pandemic Technology and the "Human Touch"

Historically, the integration of Information and Communication Technologies (ICT) in the hospitality sector focused heavily on operational efficiency, inventory management, and centralized distribution. Early advancements allowed for the aggregation of booking channels and the establishment of robust Property Management Systems (PMS) (Law, Buhalis, & Cobanoglu, 2014).

However, despite these backend efficiencies, the front-facing guest experience remained fundamentally reliant on human interaction. Subramony et al. (2017) extensively documented the critical importance of leveraging the "human touch" in service interactions, warning that over-automation risks commoditizing the guest experience, particularly in boutique and luxury segments where personalized concierge services act as a primary competitive advantage.

2.1.2 COVID-19 and the Acceleration of Contactless Systems

The onset of the COVID-19 pandemic catalyzed an unprecedented and forced acceleration in technological adoption. To ensure compliance with global health mandates and mitigate the severe socio-economic implications of the virus, small-to-medium boutique hotels rapidly pivoted toward "contactless" solutions. These reactive digital measures, ranging from QR-code menus to automated check-in kiosks, succeeded

in maintaining operational continuity but fundamentally disrupted the personalized human-centric service model (Robinson, 2024).

2.1.3 The Post-Pandemic Macro-Economic Imperative

As the industry transitions into a post-pandemic reality, the technological and economic imperatives have evolved. While the reactive contactless measures of 2020 ensured short-term survival, they created a critical service gap: high-value tourists now expect both seamless digital convenience and highly personalized concierge experiences (Robinson, 2024). Because small-to-medium boutique hotels cannot simply increase their human labor force to provide 24/7 personalized service without destroying their profit margins, they face a severe macro-economic bottleneck.

Consequently, the industry is forced to look beyond basic digitization. To successfully execute the "high-value, sustainable" recovery strategy outlined by Deloitte (2025), hotels must pivot toward advanced technological interventions capable of anticipating guest needs, personalizing itineraries, and actively driving revenue toward the local community without proportional increases in staffing costs.

However, the first wave of this transformation, basic contactless QR codes and digital check-ins, failed to deliver the personalized concierge experience required by high-value tourists, creating a critical service gap that requires advanced technological intervention.

2.1.4 Key driving factors for AI

The integration of Artificial Intelligence in the hospitality sector is not merely a technological trend but a strategic response to three core institutional and economic pressures. As boutique hotels transition away from mass tourism toward high-value, experiential travel, AI serves as the primary differentiator in the following areas:

1. Customer Experience (Hyper-Personalization)

Traditional digital tools often result in "over-automation," which risks commoditizing the guest experience. AI drives value by understanding complex guest data and preferences to restore the "human touch" through personalized digital concierge services.

2. Operational Efficiency

Small-to-medium boutique hotels frequently face a macro-economic bottleneck where they cannot increase human labor without eroding profit margins. AI acts as an autonomous medium that learns and executes tasks, such as itinerary planning and booking management, to improve operational efficiency without proportional increases in staffing costs.

3. Revenue Improvement (The Shared Value Model)

Beyond cost-cutting, AI facilitates a "Shared Business Value Model". By intelligently routing guest requests to local community partners (e.g., longtail boat captains or local markets), the AI creates new revenue streams for both the hotel and the surrounding Community-Based Tourism (CBT) network.

4. Sustainable Growth

Deep knowledge of individual customers enables the marketing team to precisely suggest products and services to the customer (Infosys, 2018).




Figure 2.1: Key Drivers for AI. *Source: Infosys (2018)*

2.2 Data Analytics in the Service Industry

Data analytics is an overarching discipline encompassing the complete management of data, including its collection, organization, storage, and analysis. While often used interchangeably with data analysis, the latter is a specific subcomponent focused on examining and arranging data sets to extract useful information.

2.2.1 What is Data Analytics?

According to BMC Software (Kidd & Hornay, 2021), data analytics is the broad field of using data and tools to make business decisions, whereas data analysis, a subset of data analytics, refers to specific actions. The primary difference between analytics and analysis is a matter of scale, as data analytics is a broader term of which data analysis is a subcomponent. Data analysis refers to the process of examining, transforming, and arranging a given data set in specific ways to study its individual parts and extract useful information.

Data analytics is an overarching science or discipline that encompasses the complete management of data. This not only includes analysis, but also data collection, organization, storage, and all the tools and techniques used.

It is the role of the data analyst to collect, analyze, and translate data into information that is accessible. By identifying trends and patterns, analysts help organizations make better business decisions. Their ability to describe, predict, and improve performance has placed them in increasingly high demand globally and across industries.

2.2.2 Hotel Data Types

From “Hotel Data Management: Solutions and Practices to Turn Information into a Valuable Asset” (AltexSoft, 2019), for the hospitality industry, the power of data helps decision-makers to solve the challenging domain-specific tasks including improving occupancy forecasting, setting competitive room prices, choosing the most profitable distribution channels, optimizing procurement operations, increasing guest loyalty, and identifying and targeting the most profitable guests (Figure 2.2; Table 2.1).




Figure 2.2: Hotel Data Map. *Source: AltexSoft (2019)***Table 2.1:** Hotel Data Types. *Source: AltexSoft (2019)*

Booking and Property Data	
Basic guest information	name, age, country of residence
Distribution channel	website, online travel agency, a metasearch engine, hotel chain, etc.
Lead time	advance time of booking
Length of stay (LOS)	duration of stay
Room price	a stay cost per night
Room history	name of guest and time per room
Key performance metrics (KPI)	average daily rate, occupancy rate, revenue per room
Service Data	
Hotel staff	number of staff working
Service time	service response times
Service KPI	performance indicators, quality of service
Service expense	expense
Service request	guest requests and complaints
Service score	guest scores for service
Room Rate Data	
Room rate	room price per day
Guest Data	
Contact information	phone number, physical and e-mail address
Demographics	age, marital status, number of children
Booking	reservation data, booking channel preferences, and past booking history
Reason	reasons for stay
Service	auxiliary services used
Food	food and beverage preferences
Cleaning	cleaning notes and housekeeping preferences
Revenue	revenue details
Payment	payment methods
Review	loyalty level (feedback)
Procurement Data	
Food	food and beverage
Furniture	furniture, fixtures and equipment
Disposable	disposables
Linen	linen and towels
Service	services
Uniform	uniforms
Print	printing

2.3 Hotel Data Analytics

The hotel industry possesses a wealth of data, encompassing a vast array of information. Unfortunately, most service providers fail to utilize this asset to its full potential. While many hoteliers collect loyalty data, they rarely delve into analytical processes to gain a deeper comprehension of their guests' behaviors, desires, and preferences. This limits their ability to develop a deeper understanding of their customers, tailor their services to meet their needs and preferences, and explore new ways to attract additional clients.

Hotel Data Analytics is often used to segment guests according to booking trends, behavior, and other factors to reveal their likelihood to respond to promotions and emerging travel trends. It is vitally important for hoteliers to be able to understand guest preferences (locations, activities, and room types), purchase behavior (frequency, length of stay, time of year), and profit potential to increase brand loyalty and wallet share of their most valuable guests.

Regrettably, financial resources are frequently allocated towards general advertising campaigns that fail to focus on specific guests or market segments with tailored promotions that are more likely to elicit a positive response. Consequently, guests may perceive that the hotel lacks concern for their individual needs, or that its services are not tailored to meet their specific requirements. Such an experience could motivate guests to switch to a competing hotel with a more personalized approach.

To truly leverage the power of analytics, hospitality organizations must distinguish between reactive and proactive decision-making. While historical data can provide reports, drilldowns, or alerts to monitor customer trends, advanced analytics can help hotels identify the root causes of these trends, forecast future outcomes, and determine the best course of action, while considering all operational constraints. By utilizing such

advanced analytics, hotels can achieve a deeper understanding of their patrons, and thereby create a significant impact on their operations (Dragosavac, n.d.).

While these internal analytics optimize hotel operations, they fail to connect the guest with external, dynamic community experiences, highlighting the need for specialized recommendation engines.

2.3.1 Customer Segmentation

Initially, the hotel's task is to identify distinctive cluster groups, followed by performing a value segmentation analysis for each group separately. Later, the hotelier can review their present business customer base and create unique clusters of business customers. For instance, they may recognize that some business customers only use the hotel for brief overnight stays, whereas others are there for longer-term events conducted at the hotel. It is feasible to break down these groups even further by examining other behavioral and relationship aspects (Dragosavac, n.d.).

2.3.2 Customer Profiling

To profile customers, a comprehensive examination of guest demographics and lifestyle traits is conducted. Characteristics such as income, family status, age, and interests in sports and culture, if available, can be added to shape a model of guests. This technique can be leveraged to generate an email roster for focused marketing campaigns aimed at both existing and potential customers. Identifying prospect profiles can be particularly advantageous in pinpointing those who are more likely to react positively to marketing and promotional proposals. Profiling can also assist in determining which market segments are the most productive and profitable (Dragosavac, n.d.).

2.3.3 Menu Engineering

A review of the sales of menu items and their contribution margins can aid in sustaining a growth of restaurant business. While menu engineering addresses menu content choices, data mining can generate reports that show the preferred menu item selections by customer segment, which can serve as a foundation for enhancing operations. For example, Applebee's has been known to utilize data mining to determine ingredient restocking amounts, leveraging a menu optimization quadrant analysis that summarizes the sales of menu items. Such analyses help the hotel decide which menu items to promote (Dragosavac, n.d.).

2.3.4 Productivity Indexing

Through a correlation of the time of order entry with the time of settlement, hoteliers can make more precise estimates of production and service times. These data offer an understanding of the average service time regarding customer turnover, along with wait time statistics. Although productivity data can be challenging to gather, this analysis provides concrete information to aid management in refining operations (Dragosavac, n.d.).

2.3.5 Customer Associations and Sequencing

Advanced big-data analytics can reveal correlations between seemingly disconnected events. For example, if a guest orders the restaurant's signature dish, they are likely to also order a small antipasto salad and a glass of Chardonnay. These paired relationships can form the basis for bundling menu items into a cohesive meal, simplifying the ordering process while ensuring customer contentment. Menu design can be adjusted to showcase such combinations as distinctive opportunities for customers. Data associations are often considered a means of persuading customers to spend more than they intended or to upsell (Dragosavac, n.d.).

2.3.6 Forecasting

Forecasting empowers restaurants to proactively exceed their clients' requirements by optimizing staffing, procurement, preparation, and menu design. Customers perceive their stay at a hotel as an experience rather than just a visit, and various activities, such as fine dining, nightly entertainment, spas, corporate seminars, or meetings, contribute to enhancing their experience. However, not all activities are equally appealing to all clients. In this regard, analytics can play a pivotal role in enabling hotels to comprehend the diverse needs of their clients more effectively (Dragosavac, n.d.)

2.3.7 Energy Consumption

Analytics can be leveraged for internal operations in the hotel industry as well. Energy consumption constitutes a significant portion of a typical hotel's utility costs, accounting for 60 to 70 percent. However, it is possible to manage costs efficiently while maintaining guest comfort by optimizing energy usage. Managers can leverage smart data to create energy profiles for their hotels. Modern software solutions can collate data from multiple sources, including weather forecasts, electricity rates, and a building's energy consumption, to develop a comprehensive "building energy profile." With the aid of a cloud-based, predictive analytics algorithm, the software can optimize whether the power supply comes from the grid or an onsite battery module (Dragosavac, n.d.).

2.3.8 The Right Room at the Right Rate

Although yield management has been a practice in the hotel industry for some time, big data offers hotels the opportunity to take revenue management to the next level by providing truly personalized prices and rooms to guests. Industry studies suggest that Marriott, a leading hotel chain, has leveraged big data analytics to predict the ideal price of its rooms to fill its hotels. By utilizing advanced revenue management algorithms that

can process data at a faster pace, Marriott has been able to combine various data sets to derive insights that are accessible to all levels, thereby enhancing decision-making capabilities (Dragosavac, n.d.).

2.4 Traditional Recommendation Systems and Their Limitations

To understand the limitations of global recommendation algorithms, one must examine the specific socio-cultural fabric of localized destinations like Koh Lanta, Krabi. Koh Lanta is not a homogenized, mass-market resort island; it is a complex demographic mosaic. Its earliest inhabitants, the Urak Lawoi (Sea Gypsies), maintain a deep, animistic connection to the marine environment. Over centuries, the island integrated Thai-Muslim fishing communities and Thai-Chinese merchants. Today, the island operates on a "Green Tourism" model, emphasizing eco-friendly practices and authentic cultural immersion, such as the Tung Yee Peng mangrove tours or the Lanta Old Town walking streets.

This Community-Based Tourism (CBT) relies on micro-enterprises. These local artisans, independent longtail boat captains, and family-run cookery schools do not possess sophisticated digital booking APIs. Consequently, standard collaborative-filtering algorithms, which rely on vast amounts of historical, digitized user data, fail entirely in this environment. This phenomenon, known as the "Cold-Start Problem," occurs because the algorithm cannot match a tourist to a local event that has no historical digital footprint, thereby trapping the tourist in generic recommendations and starving the local community of direct revenue.

2.4.1 Principles of Community-Based Tourism (CBT)

Community-Based Tourism is a sustainable framework designed to ensure that the economic benefits of travel remain within local ecosystems. Moscardo (2023) emphasizes that true sustainable tourism development requires deep community

involvement, equitable profit-sharing, and the integration of authentic local knowledge. In regions like Koh Lanta, CBT manifests through micro-enterprises: independent longtail boat operators, family-run craft markets, and localized culinary classes.

2.4.2 The Evolution from CSR to CSV

While Community-Based Tourism (CBT) establishes the social objectives for sustainable travel, realizing these goals within a capitalist framework requires a robust economic engine. Historically, the tourism industry has relied on Corporate Social Responsibility (CSR) to address community impact. CSR typically functions as a peripheral, philanthropic activity, such as a multinational hotel chain donating a fraction of its profits to local schools, while the core business model remains extractive.

In contrast, the concept of Creating Shared Value (CSV), introduced by Porter and Kramer (2011), argues that a business's economic success and social progress are inherently interconnected. CSV dictates that social impact should not be an afterthought; rather, the enterprise must generate its profit by solving a social problem. In the context of tourism, CSV requires integrating local micro-enterprises, such as longtail boat operators and local artisans, directly into the primary value chain of the hospitality sector, transforming them from passive beneficiaries of charity into active economic partners (Camilleri, 2017).

2.4.3 The Shared Value Model in Platform Economies

Applying CSV to modern digital environments reveals a significant gap in the literature. Modern tourism is dominated by massive, centralized platform economies (e.g., global Online Travel Agencies). These platforms often operate on extractive models, utilizing algorithmic popularity biases that marginalize local micro-enterprises in favor of heavily

resourced competitors, creating a systemic "algorithmic inequality" (Van Dijck et al., 2018).

A true "Shared Value Model" in a digital ecosystem must mathematically counteract this bias. It requires an architecture that explicitly values the socio-economic contributions of the community over mere historical click-through rates. By leveraging semantic vector matching rather than popularity-based ranking, an AI-driven Shared Value Model bridges the gap between the economic needs of the local vendor and the safety and relevance requirements of the tourist, ensuring equitable digital visibility for community stakeholders.

2.4.4 The "Cold-Start" Problem in Informal Micro-Economies

Despite its benefits, CBT is inherently incompatible with standard digital recommendation engines. Traditional systems primarily utilize collaborative filtering or content-based filtering, identifying patterns by analyzing vast, historical matrices of user-item interactions (Yuan & Hernandez, 2023). However, these systems suffer acutely from the "cold-start" problem.

In a CBT environment, the "items" being recommended are highly volatile. Local communities rarely maintain persistent digital APIs; their availability fluctuates based on ephemeral factors like monsoon weather or immediate supply. When a new tourist arrives for a brief two-day stay, or a local village announces a seasonal festival, traditional algorithms lack the historical data overlap required to make an accurate recommendation. Consequently, standard systems trap tourists in "filter bubbles," constantly recommending generic, globally recognized attractions (e.g., mass-market island tours) and effectively locking localized CBT communities out of the digital economy.

2.4.5 Typology of Recommendation Systems in Tourism

In the context of hospitality and e-tourism, recommendation systems are traditionally categorized into three primary architectures (Ricci, Rokach, & Shapira, 2022):

1. Collaborative Filtering (CF)

This method recommends items based on the historical preferences of similar users. While highly effective for large-scale platforms, CF suffers severely from the "cold-start" problem, rendering it incapable of recommending newly established or unrated local communities that lack sufficient historical data.

2. Content-Based Filtering (CBF)

This approach matches users to items based on static tags and historical user profiles. However, CBF struggles to adapt to the highly dynamic variables inherent in community-based tourism, such as sudden changes in local operating hours or weather-dependent cancellations.

3. Knowledge-Based Recommendation Systems (KBRS)

Unlike CF and CBF, knowledge-based systems do not rely on massive historical user logs. Instead, they recommend items based on explicit rules and deep domain knowledge, matching specific, complex user constraints against detailed item features.

To effectively mitigate the cold-start problem and process the complex constraints of modern travelers, this study frames its proposed Agentic RAG architecture as an advanced, conversational evolution of the Knowledge-Based Recommendation System (KBRS). By replacing static rule-based engines with non-parametric grounding (vector databases) and Agentic LLM orchestration, the system can dynamically evaluate user constraints against real-time community ground-truth.

2.4.6 Limitations of Traditional Recommender Systems in CBT

Traditional tourism recommender systems such as those powering major platforms like TripAdvisor or Klook predominantly rely on Collaborative Filtering (CF) or static Content-Based approaches (Ricci et al., 2022). While effective for major hotel chains with abundant historical data, these methods fail in Community-Based Tourism (CBT) environments. CBT offerings, such as seasonal fruit festivals or pop-up markets, are inherently ephemeral and lack the persistent digital footprint required for matrix factorization.

Furthermore, "Cold Start" in CBT is not a temporary phase but a permanent state; by the time a local event gathers enough ratings for a CF algorithm, it has often already concluded. Consequently, applying standard algorithms to CBT results in a "popularity bias," systematically excluding the very local, informal experiences that sustainable tourism aims to promote (Goodwin, 2020; Moscardo, 2023).

2.5 Large Language Models (LLMs) and the Hallucination Dilemma

Generative Artificial Intelligence, specifically Large Language Models (LLMs) built on the Transformer architecture (Vaswani et al., 2017), emerged as a potential solution to restore personalized digital interactions. However, deploying standard Base LLMs in the travel sector introduces severe safety and liability risks due to "algorithmic hallucinations."

Base LLMs generate text based on probabilistic word prediction drawn from their static, pre-trained parametric memory. When queried about a hyper-local, real-time environment like Koh Lanta, the LLM attempts to fulfill the prompt by mathematically guessing the most likely words, rather than verifying factual reality. As highlighted by Maynez et al. (2020), this lack of factual Faithfulness leads the AI to fabricate non-

existent restaurants, invent unsafe boat charter capacities, or confidently quote outdated prices. In the context of travel safety and international AI guidelines (OECD, 2022), these unmitigated hallucinations render generic LLMs unfit for commercial hospitality deployment (Wang et al., 2024).

2.5.1 The Rise of Generative AI in Travel

The introduction of the Transformer architecture (Vaswani et al., 2017) catalyzed a revolution in Natural Language Processing (NLP), leading to the development of powerful Large Language Models (LLMs). In the travel sector, Generative AI has demonstrated transformative capabilities in conversational fluency, language translation, and itinerary drafting (Robinson, 2024). LLMs possess the natural language capabilities to mimic the "human touch" that contactless software destroyed during the pandemic.

2.5.2 Parametric Memory and Domain-Specific Hallucinations

Despite their conversational prowess, standard Base LLMs present critical safety risks when deployed as standalone recommendation agents. LLMs generate responses based on static, pre-trained parametric memory. Because they lack access to real-time, hyperlocal data streams, they are fundamentally incapable of verifying the current truth of an ephemeral query.

This limitation leads to the phenomenon of algorithmic hallucination. Maynez et al. (2020) and Rawte et al. (2023) have extensively researched the concepts of factual "faithfulness" and LLM safety, benchmarking the harms associated with model fabrication. In the hospitality sector, these fabrications manifest dangerously: an LLM might invent a non-existent local restaurant, quote an inaccurate price for a boat charter, or recommend an activity that has permanently closed.

As established by international guidelines for trustworthy AI (European Commission, 2021; OECD, 2022), systems deployed in consumer-facing environments must demonstrate high factual reliability. Relying on ungrounded Base LLMs violates these safety mandates, as inaccurate physical recommendations can lead to financial loss or physical harm to the tourist.

2.5.3 Hallucinations in LLM-Based Systems

Hallucinations are outputs that are not supported by facts or contradict available evidence (Ji et al., 2024). They are especially problematic in domains where information changes rapidly or is location-specific, both of which characterize tourism environments. Hallucinations can manifest as fabricated entities, incorrect attributes (such as wrong prices or opening hours), temporally inconsistent claims (recommending closed events), or unsafe suggestions (Maynez et al., 2020).

Previous work shows that hallucinations increase when LLMs are asked about niche, local, or up-to-date topics that are underrepresented in training data (Kandpal et al., 2023). Grounding through retrieval can substantially reduce such errors, but the degree of improvement varies by domain. Tourism requires empirical evaluation due to its unique volatility and heterogeneity.

2.5.4 Trustworthy AI

Trustworthy AI frameworks, including those by OECD and the European Commission, emphasize reliability, robustness, transparency, and safety (European Commission, 2021; OECD, 2022). In generative AI, two practical evaluation metrics have emerged: faithfulness, which measures alignment between generated content and reference data, and maliciousness, which measures the presence of harmful, unsafe, or inappropriate content (Wang et al., 2023; Rawte et al., 2023). These metrics have been applied in

safety-critical domains such as healthcare and legal reasoning, but tourism research rarely adopts them, despite the potential physical and social impact of inaccurate or unsafe recommendations.

2.6 Community Data Governance in AI Ecosystems

The transition toward an AI-driven Shared Value Model necessitates a profound shift in how empirical data is collected, stored, and governed. Traditional AI pipelines rely on massive, centralized data harvesting, often commodifying user data without reciprocal benefit to the data creators (Zuboff, 2019). When deploying Generative AI in vulnerable CBT environments like Koh Lanta, extractive data practices risk exploiting local knowledge, such as proprietary cultural practices or hyperlocal weather strategies, for corporate algorithmic training without compensating the community.

Community Data Governance presents a localized, decentralized alternative. Drawing upon Ostrom's (1990) principles of governing the commons, modern data governance in localized AI treats community data not as a raw resource to be extracted, but as a shared asset to be stewarded. In a Tripartite knowledge-sharing ecosystem, data governance requires strict access controls, ethical sovereignty, and clear reciprocal value.

When local communities input highly volatile data, such as real-time physical safety constraints, dynamic pricing, or operational availability, into the hotel's NoSQL database, the governance framework ensures this data is utilized exclusively to route economic traffic back to those specific vendors via the Agentic RAG architecture. By explicitly embedding community data sovereignty into the system's architecture, the platform complies with regional data protection frameworks like Thailand's PDPA (Personal Data Protection Committee, 2019) while mathematically ensuring that the AI operates as a tool for local economic empowerment rather than extraction (Micheli et al., 2020).

2.7 Retrieval-Augmented Generation (RAG) and Agentic Systems

To neutralize the hallucination dilemma, this research leverages Retrieval-Augmented Generation (RAG) integrated with an Agentic workflow. Unlike standard LLMs, an Agentic AI acts as an autonomous planner (Wang et al., 2024). When a tourist submits a request, the Agentic system pauses the natural language generation process. It first triggers a retrieval pipeline, in this study, powered by the Firebase Genkit framework, to query a centralized, real-time vector database of verified communities.

Once the current, factual context is retrieved (e.g., verifying that a specific longtail boat operator is available today and costs exactly 1,500 THB), the RAG architecture forces the Gemini LLM to generate its response strictly using the retrieved data. By subordinating the AI's parametric memory to the authoritative local database, RAG effectively bridges the gap between conversational fluency and verifiable factual grounding, solving the Cold-Start problem and ensuring the economic value of the recommendation remains within the localized CBT network.

2.7.1 The Agentic Paradigm

To resolve the passivity and unreliability of Base LLMs, recent advancements have shifted focus toward Agentic AI (Wang et al., 2024). Unlike standard conversational models that simply react to single prompts, AI agents are endowed with autonomous reasoning and planning workflows. They can break down complex user constraints (e.g., budget, group size, and dietary restrictions) into sequential tasks, actively seeking external data to inform goal-oriented decisions (Gill, 2025).

2.7.2 Retrieval-Augmented Generation (RAG)

RAG architectures combine external knowledge retrieval with generative reasoning (Lewis et al., 2020; Nogueira et al., 2024). Queries are encoded into embeddings, matched against document vectors, and the retrieved passages are fed into the LLM as context. Dense retrieval models such as DPR, ColBERT, and related architectures have demonstrated superior semantic matching performance compared with sparse methods like BM25 (Khattab & Zaharia, 2020; Karpukhin et al., 2020).

RAG has achieved strong results in knowledge-intensive tasks, such as open-domain question answering, legal document reasoning, and clinical decision support (Chen et al., 2024; Xiao & Wu, 2023). Nevertheless, tourism applications remain limited. Only a small number of studies explore neural retrieval for tourism search (Yuan & Hernandez, 2023), and there is no established benchmark for evaluating RAG against base LLMs in CBT contexts. This study contributes to filling that gap.

2.7.3 The Mechanics of RAG

Retrieval-Augmented Generation (RAG), formally introduced by Lewis et al. (2020) and Karpukhin et al. (2020) for open-domain question answering, represents the architectural bridge between fluid text generation and verifiable factual data. RAG fundamentally alters the AI's workflow: before generating a response, the system queries an external, authoritative knowledge base (such as a centralized vector database), retrieves relevant context, and forces the LLM to strictly ground its generated answer within those retrieved facts (Sivan, 2024).

Recent studies have begun exploring RAG in tourism. Banerjee, Satish, and Wörndl (2025) demonstrated its efficacy in enhancing sustainable city trips by neutralizing parametric hallucinations. By treating real-time local data as the absolute ground truth, RAG ensures that the AI's recommendations are faithfully aligned with physical reality.

2.7.4 AI Agents and Autonomous Decision Systems

AI agents have evolved from scripted rule-based chatbots to systems capable of intent parsing, multi-step planning, and tool invocation (Shinn et al., 2023; Xi et al., 2025). Modern agentic frameworks allow LLMs to decide when to call tools such as retrievers, search APIs, or function calls, following patterns such as chain-of-thought, planning, and reflection. These advances have improved performance in complex reasoning tasks and task-oriented dialogue.

Despite these developments, most tourism applications still use relatively simple chatbots or static Q&A systems that rely solely on the LLM's parametric memory. By failing to incorporate retrieval or agentic planning, standard systems are fundamentally unequipped to process dynamic, community-level variables. This gap motivates the integration of RAG into agentic tourism systems.

2.8 Identification of the Research Gap

A synthesis of the current literature reveals a clear technological progression, but also a glaring gap. While extensive research exists on the socio-economic benefits of CBT (Moscardo, 2023), and separate computer science literature thoroughly documents the efficacy of RAG in enterprise applications (Lewis et al., 2020), there is a profound lack of empirical research bridging these two domains.

Currently, RAG architectures are predominantly built for corporate data silos. There is no existing, validated architectural blueprint that applies Agentic RAG to the highly unstructured, informal micro-economies of localized tourism in Southeast Asia. Furthermore, while theoretical frameworks suggest AI can aid sustainability, there is a lack of empirical measurement regarding how effectively an Agentic system can

eradicate the specific taxonomies of hallucination (fabrication, temporal decay, and attribute misalignment) when applied to real-world boutique hotel operations.

This dissertation directly addresses this research gap. By designing, implementing, and evaluating a tripartite "Hotel, Community, and Tourist Relationship Platform" utilizing Firebase Genkit and RAG, this study transitions Agentic AI from a theoretical corporate tool into a validated, shared-value engine for sustainable Community-Based Tourism.

2.9 Firebase Genkit as an AI Framework

This research utilizes the Firebase Genkit platform to develop the recommendation agent workflow. Genkit is an open-source TypeScript toolkit designed to simplify the integration of AI-powered features into applications (Firebase, n.d.). Serving as a unified interface, it enables the integration of AI models from various providers and offers capabilities like Structured Data Generation and Tool Calling. In this study, Genkit was essential for orchestrating the interaction between the Firestore database retrieval and the LLM generation.

While the theoretical benefits of Retrieval-Augmented Generation are well-documented, operationalizing these architectures requires specialized orchestration tools. This research utilizes the Firebase Genkit platform to systematically develop and manage the recommendation agent workflow. Genkit is an advanced, open-source TypeScript toolkit explicitly designed to abstract the complexities of integrating AI-powered features, such as autonomous agents, workflows, and semantic recommendation engines, into production environments.

In the context of this study, Genkit was essential for moving beyond static API calls to create a robust, autonomous Agentic RAG framework. Specifically, it provided the structural capacity to:

- **Define and Execute Complex Flows**

Genkit allowed for the precise architectural design of execution flows, successfully orchestrating the intricate interactions between the retriever (querying the NoSQL database) and the generator (prompting the Large Language Model).

- **Integrate with Cloud Infrastructure**

It facilitated seamless integration with Google Cloud services, bridging Firestore for real-time data storage with Google's Gemini models to ensure an uninterrupted, low-latency data flow.

- **Implement Custom Retrieval Logic**

Rather than relying on generic search algorithms, Genkit enabled the implementation of custom semantic retrieval logic that was mathematically optimized for the specific data structures of the Koh Lanta community knowledge base.

- **Automate Performance Evaluation**

Crucially, it allowed for the integration of continuous evaluation protocols, utilizing automated metrics to assess recommendation quality directly within the development pipeline.

2.9.1 Recommendation Evaluation Metrics

To rigorously assess the quality, accuracy, and safety of the recommendation agent, the system employs standardized metrics derived from the Genkit Evaluation tool and the Ragas framework. Evaluating Agentic RAG systems requires advanced methodologies to ensure both semantic relevance and strict adherence to retrieved facts. The following evaluation frameworks were utilized:

1. Maliciousness

This assessment uses LLM evaluators for submissions based on predefined (e.g., correctness, harmfulness via Ragas Critiques) or custom aspects (Ragas Team, n.d.).

The "strictness" parameter aids self-consistency (Ragas Team, n.d.). Calculation involves LLM calls using defined criteria:

Step 1: Critic definition prompts LLM multiple times, asking if submission is harmful (e.g., "Does the submission cause... harm...?"). Three verdicts were collected.

Step 2: Majority vote determines the binary output (Yes/No).

2. Answer Relevancy

Measures how pertinent the generated answer (recommendation) is to the user's query (Ragas Team, n.d.). A high score indicates effective understanding. Purpose is ensuring useful, targeted information, essential for user satisfaction (Ragas Team, n.d.). Assessed by comparing recommendations to the original prompt:

- **Manual Evaluation**

Humans rate relevancy on a scale (e.g., 1-5); average gives score.

- **Automated Evaluation**

Uses NLP for semantic similarity via methods like Embedding Similarity (e.g., Sentence-BERT cosine similarity), Keyword Matching, or Topic Modeling (e.g., LDA).

Score is numerical/categorical; higher means more relevant. Crucial for perceived usefulness and trust (Ragas Team, n.d.). A high score indicates the ability to provide tailored recommendations that meet needs. Defined as the mean cosine similarity of the original question embedding (Ragas Team, n.d.).

$$answer\ relevancy = \frac{1}{N} \sum_{i=1}^N \cosinesimilarity(E_{gi}, E_o)$$

$$answer\ relevancy = \frac{1}{N} \sum_{i=1}^N \frac{E_{gi} \cdot E_o}{\|E_{gi}\| \|E_o\|}$$

Where:

- E_{gi} is the embedding of the generated question.
- E_o is the embedding of the original question.
- N is the number of generated questions, which is 3 by default.

3. Faithfulness

Measures factual consistency of the generated answer against the provided context in RAG systems, ensuring recommendations are grounded (Ragas Team, n.d.). Verifies accuracy and reliability, crucial for preventing misinformation and maintaining trust (Ragas Team, n.d.). Assessed by comparing generated recommendations to retrieved context:

- **Manual Evaluation**

Humans review context and recommendation, assess statement support, and assign a score.

- **Automated Evaluation**

Uses NLP to check information presence via methods like Text Overlap, Semantic Similarity, or Fact Verification.

Score is numerical/categorical; higher means more factually consistent (Ragas Team, n.d.). High Faithfulness demonstrates a system that generates recommendations based on real information, enhancing user confidence.

$$Faithfulness\ score = \frac{|Number\ of\ claims\ in\ the\ response\ supported\ by\ the\ retrieved\ context|}{|Total\ number\ of\ claims\ in\ the\ response|}$$

2.10 The "Hotel, Community, and Tourist Relationship Platform"

Building upon the socio-economic need for decentralized tourism and the technological capabilities of Retrieval-Augmented Generation (RAG), this research proposes a tripartite conceptual model. This model fundamentally redesigns the information flow within a destination by positioning the boutique hotel as an intelligent Knowledge Hub.

1. The Tourist (Demand Generation)

Modern tourists, particularly in the post-COVID era, demand hyper-personalized, authentic, and safe experiences. However, they lack the local knowledge required to navigate informal, unstructured community markets safely. In this model, the tourist provides complex, multi-variable constraints (e.g., length of stay, physical mobility, group size) via natural language to the hotel's AI interface.

2. The Hotel as the Knowledge Hub (Agentic Orchestration)

At the center of the model is the boutique hotel. Rather than relying on human concierge staff, which is economically unfeasible to scale 24/7, the hotel utilizes an Agentic RAG architecture. The hotel's role is to act as a secure intermediary. It ingests the tourist's demand, filters it for safety and viability, and matches it against verified local data. The hotel becomes the curator of trust, protecting both the tourist from unsafe experiences and the local communities from miscommunications.

3. The Local Community (Decentralized Supply)

The final node consists of the local Koh Lanta communities (e.g., Community-Based Tourism enterprises, independent boat captains, local markets). In traditional models, these micro-entrepreneurs are invisible to global booking platforms. In this conceptual model, they directly supply the "ground truth" data (schedules, pricing, and safety limits) into the hotel's NoSQL database. By conceptualizing the hotel as a technological and relational hub, this model ensures that the implementation of

Generative AI does not isolate the tourist, but rather embeds them safely and economically into the local ecosystem.

2.11 Summary of Chapter 2

This chapter has systematically reviewed the multi-disciplinary literature required to understand and solve the modernization challenges within Thailand's post-pandemic hospitality sector. The review began by examining the macro-economic shift toward high-value, experiential travel, identifying a critical need to artificially restore the personalized "human touch" that was lost during the rapid adoption of reactive, contactless technologies. By analyzing the socio-economic realities of localized Community-Based Tourism (CBT) in Koh Lanta, this chapter demonstrated that traditional, collaborative-filtering recommendation systems are fundamentally unequipped to handle the dynamic nature of micro-enterprises, failing primarily due to the "cold-start" problem.

While the advent of Generative Artificial Intelligence and Large Language Models (LLMs) presented a theoretical solution to this hyper-personalization gap, the literature firmly establishes that ungrounded Base LLMs pose severe safety and liability risks. Because standard models rely purely on static parametric memory, they are prone to dangerous algorithmic hallucinations, fabricating local entities, quoting inaccurate prices, and violating international guidelines for Trustworthy AI.

Consequently, this review identified Agentic Retrieval-Augmented Generation (RAG) as the necessary architectural bridge. By utilizing an Agentic AI framework, such as Firebase Genkit, the system is forced to bypass its static memory and retrieve real-time, authoritative data directly from the local community before generating a response. Ultimately, the literature reveals a clear research gap: there is currently no validated,

operational blueprint for applying Agentic RAG to synchronize hotel databases, tourist personas, and unstructured community micro-economies. Addressing this exact technological and socio-economic gap forms the methodological basis of this dissertation, the architecture of which is detailed in Chapter 3.

CHAPTER 3 METHODOLOGY AND SYSTEM ARCHITECTURE

This chapter details the research methodology, system design, and technical architecture employed to address the limitations of traditional recommendation systems in Community-Based Tourism (CBT). Rooted in a constructive research paradigm, this study seeks to design, implement, and validate a novel artifact, an Agentic Retrieval-Augmented Generation (RAG) platform.

This chapter details the methodology employed to design, develop, and thoroughly evaluate the proposed Tripartite Ecosystem and its underlying Agentic RAG architecture. To ensure a comprehensive assessment of the system, the research design is divided into both technical development and rigorous mixed-methods validation. The chapter outlines the data collection and pre-processing methods, followed by the system architecture and database schema design.

The evaluation of the system is then conducted in two distinct phases: first, through a quantitative baseline analysis using automated metrics for safety and factual grounding; and second, through a qualitative evaluation outlined in section Human Expert Review Setup, which details the framework for benchmarking the AI's recommendations against actual local hospitality professionals to ensure cultural accuracy and practical viability.

3.1 Research Design and Paradigm

This research employs an applied, constructive methodology to investigate the efficacy of Agentic AI and Retrieval-Augmented Generation (RAG) in enhancing tourist recommendation systems for small-to-medium boutique hotels. Constructive research is a widely accepted paradigm in applied computer science, focusing on solving practical,

real-world problems through the creation of innovative artifacts, in this case, an integrated digital ecosystem.

The methodology encompasses three core phases: the development of a cloud-based relationship platform, the generation of a synthetic dataset based on real-world historical contexts, and the rigorous quantitative evaluation of the AI recommendation agent. The contextual setting for this research is Koh Lanta, Krabi. Because sustainable tourism development requires deep integration of local knowledge (Moscardo, 2023), the methodology combines qualitative, domain-specific insights from local communities to construct the foundational knowledge base, followed by a quantitative evaluation of the AI's performance utilizing standardized benchmarking.

3.1.1 Data Governance and PDPA Compliance Framework

The deployment of the Tripartite knowledge-sharing ecosystem within the platform necessitates a rigorous data governance framework to ensure legal compliance with Thailand's Personal Data Protection Act (PDPA) (Personal Data Protection Committee, 2019). The architecture establishes strict boundaries regarding data ownership, management, and processing. Within this ecosystem, the participating Boutique Hotel assumes the legal role of the Data Controller, holding authority over the collection and purpose of guest data. The Agentic RAG platform, operating via a secure cloud infrastructure, functions as the Data Processor. The tourists and local vendors retain absolute sovereignty over their profiles as the Data Subjects.

To mathematically guarantee PDPA compliance, the Agentic RAG architecture is engineered around the principle of Data Minimization. The AI orchestrator strictly extracts only the contextual constraints required to execute a semantic vector match without retaining a permanent, identifiable behavioral log. Furthermore, the Firebase

NoSQL document structure inherently supports the "Right to be Forgotten"; the deletion of a Data Subject's core document automatically removes their data points from the active retrieval ecosystem, ensuring the generative AI cannot recall or hallucinate their information.

3.2 System Architecture: The Tripartite Ecosystem

To transition the Large Language Model from a generic text generator into a localized Agentic AI, strict system constraints were engineered into the architecture. These constraints act as the guardrails for semantic vector matching, ensuring the AI cannot generate recommendations that violate the physical safety or operational reality of the Koh Lanta community. The constraints are categorized into three primary vectors.

1. Physical and Environmental Constraints

Real-time data regarding monsoon weather warnings, maritime safety statuses for longtail boats, and physical accessibility requirements for tourists.

2. Operational Constraints

Dynamic availability of local micro-enterprises, including daily opening hours, real-time pricing fluctuations, and capacity limits.

3. Hyper-Personalized Constraints

Specific user-defined requirements, such as strict dietary restrictions or preferences for authentic regional cuisine. By enforcing these constraints, the architecture strictly limits generation to mathematically validated local data, effectively mitigating harmful hallucinations.

The architectural framework of this study is designed as a Tripartite Application Ecosystem. Rather than functioning as siloed products, these three portals act as synchronized data-ingestion nodes, engineered to overcome the fragmented digital

landscape of small-to-medium hospitality enterprises. The ecosystem captures unstructured and structured inputs, routing them into a unified NoSQL database to feed the generative AI pipeline.

- **Hotel Front**

Functioning as the primary deterministic data node, this web interface handles staff inputs such as guest booking IDs, stay durations, and room assignments. It writes directly to the central database, establishing the baseline parameters and constraints required by the AI's Context Retriever.

- **Hotel Community**

Operating as an unstructured data collection tool, this interface allows local communities to dynamically update their offerings (e.g., current prices, operating hours, menus). This continuous data ingestion is crucial for mitigating the Temporal Decay hallucinations identified in Chapter 2, ensuring the vector database reflects real-time community micro-economies.

- **Hotel Guest**

Serving as the direct end-user endpoint, this mobile application captures explicit tourist queries and implicit persona data. It acts as the trigger mechanism, taking the user's prompt, executing the semantic search across the unified database, and delivering the Agentic AI's finalized output to the user's screen.




Figure 3.1: Architectural of the Application and the Recommendation Agent. *Source: Author (2025)*

3.2.1 The Context Retriever and Real-Time Data Ingestion

In a standard Agentic RAG pipeline, the retrieval mechanism strictly prioritizes semantic similarity (Cosine Distance) to match the user's query with the database documents. However, relying solely on semantic proximity in a Community-Based Tourism ecosystem risks reinforcing algorithmic inequality, as highly optimized profiles from larger entities may consistently dominate the top retrieval results. To mathematically counteract this bias, the proposed architecture integrates an application of the Sustainability-Augmented Reranking (SAR) framework (Banerjee et al., 2024). Following the initial vector retrieval, the SAR mechanism is applied as a secondary weighting function. This function introduces a socio-economic modifier to the baseline semantic score. Within the context of the Koh Lanta Shared Value Model, the reranking algorithm slightly elevates the retrieval priority of local micro-enterprises, provided they meet strict safety constraints, ensuring the final generative output actively distributes economic visibility.

To effectively mitigate the hallucinations and temporal decay inherent in ungrounded Large Language Models (Maynez et al., 2020), the proposed architecture fundamentally subordinates parametric generation to real-time data retrieval. The core mechanism enabling this is the `contextRetriever`, a specialized backend module engineered within the Firebase Genkit framework.

Rather than relying on static datasets, the `contextRetriever` continuously synchronizes with a centralized NoSQL Firestore vector database. As hotel staff, guests, and local communities interact with their respective interfaces within the Tripartite Ecosystem, their inputs are immediately ingested, embedded, and stored. To overcome the traditional "cold-start" problem (Yuan & Hernandez, 2023) and construct a highly accurate, grounded context for the Agentic planner, the retriever actively queries across several dynamic data collections:

- **Orders and Baseline Activity Collection**

The system filters historical booking data (via `book_id`) matching the specific `hotel_id`. This establishes a baseline of guest behavior, stay duration, and foundational demographic constraints before a query is even made.

- **Implicit Interest (Views and Reviews)**

To capture the nuanced, evolving preferences of modern travelers, the retriever ingests explicit ratings and implicitly viewed items. As defined by the system's engineered prompt rules, products a user has engaged with or looked at, even without completing a transaction, serve as critical indicators of potential interest.

- **Dynamic Partner Inventory**

The system aggregates real-time service offerings directly from local Community-Based Tourism (CBT) partners. It parses highly specific, volatile attributes, including

item descriptions, strict physical capacities (e.g., distinguishing between a "Private Charter Longtail Boat" for 1-4 people versus 7-8 people), current pricing, and live community review scores.

By continuously extracting data from these distinct but interconnected streams, the contextRetriever ensures that the subsequent generative phase is constrained by the absolute, real-time "ground truth" of the Koh Lanta community. This architecture guarantees that the Agentic AI only recommends services that physically exist and are currently available, thereby neutralizing the risks of algorithmic fabrication.

3.2.2 The RecommendDotPrompt and Generation Phase

Once the contextRetriever successfully compiles the localized facts, the data is passed to the recommendDotprompt module. The Gemini 2.0 Flash model is instructed to act as an autonomous agent. Rather than relying on its pre-trained global memory, which leads to hallucination (Maynez et al., 2020), the Agentic AI cross-references the tourist's profile against the freshly retrieved community inventory, thereby ensuring mathematical and contextual accuracy in its final recommendation.

When a guest initiates a request via the Mobile Guest App, the system performs a semantic search within the vector database. It retrieves the top-k contextually similar chunks of data based on the user's query. These retrieved, real-time localized contexts are then passed to the Gemini 2.0 Flash model via Genkit's prompt management system. By cross-referencing the tourist's profile against the freshly retrieved community inventory, the Agentic AI generates highly personalized, constraint-specific itineraries grounded entirely in factual, real-time data.

3.3 Agentic RAG Implementation via Firebase Genkit

The selection of the foundational LLM within the Agentic RAG architecture required balancing latency, tool-calling capabilities, and seamless integration with the cloud infrastructure. Gemini 2.0 Flash was selected as the orchestrating engine. While larger, more parameter-dense models exist, they often introduce significant latency overhead that degrades the real-time user experience in mobile travel applications. Gemini 2.0 Flash was explicitly designed for high-speed, multimodal reasoning and exhibits state-of-the-art proficiency in function-calling (tool-use), a critical requirement for an Agentic architecture utilizing the Firebase Genkit to autonomously query NoSQL databases.

The core computational engine of the proposed system is the Recommendation Agent, a cloud API service built using Firebase Genkit. Genkit is an open-source framework designed to simplify the integration of AI-powered features into applications, serving as the unified interface to orchestrate the interaction between Google Cloud services, the vector database, and the Large Language Models (LLMs) (Firebase, n.d.).

3.3.1 Database Schema and Prompt Engineering

To manage the complex data flow, the recommendation agent performs an embedding search within the vector database across five distinct tables: hotel services, community services, guest activities, historical orders, and service reviews. When the system retrieves the top-k contextually similar chunks, they are routed to the LLM. To strictly control the model's tone, format, and generative length, the researcher utilized Genkit's Dotprompt tool. Dotprompt allows prompts to be managed as executable code, separating the logic from the application code, with iterations continuously tested and refined via the Genkit Developer UI. Both the Base LLM and the RAG Agent were standardized using Gemini 2.0 Flash as the underlying computational model to ensure a consistent experimental baseline (see Appendix C).

3.4 System Evaluation Procedures

Evaluating the efficacy of Generative AI within specialized, high-risk domains like community-based tourism requires a rigorous, multi-layered approach. Relying solely on automated metrics can obscure cultural nuances, while relying solely on human evaluation lacks scalability (Zheng et al., 2023). Therefore, this research employed a hybrid evaluation methodology, divided into a programmatic automated phase and a subsequent qualitative manual phase, to ensure both mathematical rigor and real-world applicability.

3.4.1 Experimental Setup

The 150 queries were processed simultaneously through both systems. The Base LLM configuration received only the persona prompt, relying entirely on its parametric memory. The RAG Agent received the persona prompt alongside the retrieved documents from the hotel and community datasets. To eliminate bias during evaluation, all outputs were anonymized and randomized, ensuring that human annotators did not know whether a given answer came from the base model or the RAG agent.

3.4.2 Data Synthesis and Persona Generation

Due to the safety constraints and privacy regulations associated with deploying experimental AI systems directly to live consumers (European Commission, 2021), a pre-deployment validation methodology was adopted. However, to ensure the simulated environment accurately reflected real-world hospitality dynamics, the baseline attributes of the synthetic personas were empirically grounded in historical customer segmentation modeling.

3.4.3 Justification for Synthetic Data Utilization

To rigorously evaluate the proposed Agentic Retrieval-Augmented Generation (RAG) architecture, this study utilized a dataset of 50 synthetically constructed traveler personas rather than historical, real-world hotel booking logs. The decision to employ synthetic data over empirical user data was driven by four methodological, ethical, and regulatory imperatives:

- **Ethical Compliance and AI Safety**

Generative AI models, particularly zero-shot baseline Large Language Models (LLMs), are prone to hallucinations that can result in unsafe recommendations. In the context of maritime and community-based tourism, deploying an unverified recommendation engine to live tourists presents severe liability risks, such as suggesting boat tours during monsoon warnings or routing tourists to non-existent, unsafe locations. Aligning with international Trustworthy AI frameworks (e.g., OECD, 2022), the use of synthetic personas provided a necessary, closed "sandbox" environment. This allowed the research to stress-test the system's safety guardrails and achieve a 0.00 Maliciousness score prior to exposing human subjects to potential physical or financial harm.

- **Regulatory Compliance and Data Privacy (PDPA)**

Utilizing empirical query data requires the ingestion of actual hotel guest records, travel itineraries, and demographic profiles. Under Thailand's Personal Data Protection Act (PDPA), processing and analyzing such Personally Identifiable Information (PII) for experimental AI research without explicit, multi-tiered consent constitutes a regulatory violation. The construction of synthetic data entirely circumvents PII constraints, enabling large-scale, automated testing without compromising guest privacy or breaching legal data handling protocols.

- **Experimental Control and Edge-Case Engineering**

Historical booking data is often statistically biased toward the "average" tourist, which is insufficient for testing the limits of an AI reasoning engine. A core objective of this study was to solve the "cold-start" problem by matching highly constrained tourists with niche Community-Based Tourism (CBT). Synthetic data generation allowed the researcher to deliberately engineer complex "edge-cases"—such as conflicting physical capacity limits, restricted temporal windows (e.g., 2-day stays), and specific ethical preferences. This intentional stress-testing mathematically proved the system's ability to reason through complex constraints rather than merely retrieving generic, median data.

- **Baseline Reproducibility for Programmatic Evaluation**

To scientifically benchmark the performance of the proposed Agentic RAG system against the Base LLM, the evaluation required a deterministic and perfectly controlled environment. The automated Ragas (Retrieval Augmented Generation Assessment) framework necessitates that both AI systems process the exact same inputs under identical conditions to calculate comparative metrics (Faithfulness and Answer Relevancy). Synthetic personas provided a stable, reproducible dataset that eliminated the external noise and variability inherent in live user queries, ensuring the statistical validity of the final performance improvement metric.

3.4.4 Dataset Preparation and Chunking

Prior to executing the simulated queries, the foundational knowledge base was systematically processed and formatted to ensure optimal semantic retrieval by the Agentic RAG system. The raw, unstructured data were categorized into three primary knowledge corpora (see Appendix A):

- **Hotel Operational Dataset**

This corpus encompassed internal facility parameters, including room categorizations, dynamic dining menus, on-site amenities, and real-time service capacities.

- **Community-Based Tourism (CBT) Dataset**

Serving as the hyper-local ground truth, this dataset captured the highly volatile variables of the local micro-economy, including vendor schedules, seasonal festival dates, exact pricing, and the strict physical safety capacities of independent operators (e.g., longtail boat limits).

- **Guest Experience Dataset**

This corpus compiled anonymized historical FAQs and implicit user preferences, allowing the semantic engine to better map colloquial tourist queries to specific local events.

In Retrieval-Augmented Generation architectures, the granularity of data segmentation directly impacts the precision of the semantic search. To optimize this, the combined data were programmatically segmented into discrete chunks of 250 to 350 tokens. This specific token window was strategically selected to balance contextual completeness with retrieval efficiency; it ensures that each chunk contains a complete, self-contained thought or constraint without overwhelming the LLM's processing context.

Once segmented, these chunks were converted into dense vector embeddings and stored within the NoSQL database. During the active evaluation phase, the retrieval mechanism was calibrated to a Top-k = 5 setting. By restricting the semantic search to the five most highly correlated data chunks, the system effectively bounds the Agentic Planner's reasoning, providing it with enough verified context to formulate a personalized itinerary while mathematically restricting its ability to hallucinate unverified information.

3.4.5 Persona Construction and Diversity

The foundational dataset consists of 50 synthetic traveler personas (Salminen et al., 2018). This sample size was derived via a Combinatorial Constraint Matrix designed to stress-test the system's safety guardrails and achieve theoretical saturation (Saunders et al., 2018). The matrix multiplies key variables: 5 cultural origins $\times$ 5 physical/dietary constraint levels $\times$ 2 economic profiles. This yields 50 distinct traveler archetypes. Demographic baselines were aligned with empirical inbound tourism data to prevent algorithmic representation bias (Ministry of Tourism and Sports, 2023).

To evaluate the AI's contextual adaptability, each of the 50 personas was subjected to Temporal Prompting based on the canonical stages of the customer journey (Lemon & Verhoef, 2016). Each persona generated 3 distinct queries representing different temporal stages (pre-arrival planning, immediate on-site activity, and constraint-shifted alternative). This multiplier (50×3) resulted in exactly 150 complex queries fed into the Agentic RAG pipeline, generating 150 highly localized activities for evaluation.

To establish the empirical blueprint for the simulated test environment, an unsupervised K-Means clustering algorithm was applied to the dataset of persona queries. The feature set for the clustering model was constructed using MultiLabel Binarization of the synthesized tourism activities (e.g., "Explore Local Markets," "Take a Cooking Class"). To ensure the algorithm grouped the personas based on contextual behavioral intent rather than basic keyword frequency, Cosine Distance was utilized as the primary distance metric.

The optimal number of clusters (k) was determined iteratively using Silhouette Analysis. The Silhouette Score peaked at 0.523 for $k = 3$, mathematically validating the existence of three distinct, stable customer archetypes within the dataset. The empirical outputs of

these cluster centroids and their associated mean cosine distances are fully detailed in Section 4.1.

These personas (stored in `user_profile.json`) were generated using a seed prompt that randomized key attributes including age, budget (budget vs. luxury), and travel goals to ensure a broad spectrum of tourist intentions. Each persona included a profile summary, explicit goals, and relevant constraints. Specific engineered parameters included:

- **Nationality and Cultural Context**

Ranging from domestic tourists (e.g., Anong Suksamran, Thailand) to international travelers (e.g., David Miller, USA; Kenji Tanaka, Japan).

- **Length of Stay**

Varying from short-term layovers (1-2 days) to extended vacations (up to 14 days). This metric is critical, as short stays typically trigger the highest failure rates in traditional collaborative-filtering models due to the cold-start dilemma.

- **Activity Preferences**

Mapped via the `user_activity.json` dataset, personas were assigned overlapping but distinct goals, such as "Explore the Islands by Longtail Boat," "Take a Cooking Class," or "Visit a Museum."

By feeding these varied configurations into the system as simulated inputs (via `user_order.json` files), the research artificially replicates the chaotic, real-time demand environment of a fully booked boutique hotel in Koh Lanta (see Appendix E).

3.5 Automated Evaluation via Firebase Genkit and Ragas

The generated outputs were evaluated using the automated Ragas framework. It is critical to contextualize these metrics within the domain of Community-Based Tourism.

Standard Ragas benchmarks often cite a 0.70 threshold for "high quality" output; however, this standard is primarily derived from highly structured, static, and encyclopedic datasets. Koh Lanta's CBT ecosystem is fundamentally volatile and unstructured. The success metric in this hyperlocal context is mathematically proving the architectural leap from an ungrounded baseline. While a zero-shot Base LLM severely struggled with this unstructured community data, the proposed Agentic RAG architecture achieved a score of 0.56. This relative increase represents a significant architectural validation, establishing a new operational baseline for localized semantic search.

The first phase of the evaluation was designed to test the Agentic RAG system at scale against a baseline, zero-shot Large Language Model. To execute this, the study leveraged Firebase Genkit's built-in evaluation tools, which allow for the programmatic routing of synthetic datasets through the AI flows. The 150 simulated queries (generated from the 50 tourist personas) were processed automatically, capturing the inputs, retrieved contexts, and final generated responses.

To mathematically score these outputs, the evaluation utilized the Ragas (Retrieval Augmented Generation Assessment) framework (Es et al., 2024). Ragas operates on an "LLM-as-a-judge" paradigm, which has been shown to effectively approximate human scoring in natural language generation tasks while eliminating manual fatigue and bias (Zheng et al., 2023). Three specific quantitative metrics were calculated for every output:

The anonymized outputs were quantitatively analyzed using the standardized Ragas evaluation framework (Ragas Team, n.d.). Five trained annotators independently rated the outputs, with any disagreements resolved by a majority vote. The study focused on three critical metrics:

- **Faithfulness**

This metric assesses alignment with retrieved evidence (or ground truth for the Base LLM). High faithfulness ensures the AI is not fabricating non-existent communities or outdated prices.

- **Answer Relevancy**

This measures semantic fulfillment, the degree to which the answer addresses user intent and aligns with specific constraints (length of stay, group size, activity preference).

- **Maliciousness**

A critical safety benchmark evaluating the presence of harmful, toxic, or physically unsafe recommendations.

Faithfulness and Relevancy were measured on a continuous scale of 0 to 1, where scores > 0.7 were classified as "High Quality". Maliciousness was evaluated as a strict binary flag; any detection of toxic, unsafe, or culturally insensitive content resulted in an immediate failure state (score of 1.00) (see Appendix B).

3.6 Manual Validation via Human Domain Experts

While automated frameworks excel at calculating algorithmic faithfulness, they lack the intrinsic domain knowledge required to judge the absolute feasibility of a localized itinerary. Furthermore, as established in foundational hospitality literature, the evaluation of service quality must ultimately account for the "human touch" and cultural alignment (Subramony et al., 2017). To corroborate the Ragas scores, a second phase of manual evaluation was conducted using a panel of human domain experts.

This panel consisted of local boutique hotel managers and community tourism operators familiar with the precise logistical constraints of Koh Lanta. The human experts were

presented with blinded, side-by-side case studies containing the tourist persona, the query, the Base LLM response, and the Agentic RAG response. The experts evaluated the outputs based on three qualitative criteria:

- **Operational Feasibility**

Can this itinerary actually be executed in Koh Lanta today? (e.g., verifying that the travel time between two recommended locations is physically possible).

- **Safety and Temporal Accuracy**

Did the system correctly interpret seasonal and weather-related constraints to protect the guest?

- **Community Alignment**

Does the recommendation accurately reflect the capabilities and offerings of the local communities without overpromising?

By triangulating the automated Genkit metrics with the qualitative assessments of local human experts, the methodology ensures that the evaluated system is not merely technically sound, but operationally viable as a "Silicon Concierge" for the Koh Lanta community.

3.6.1 Panel Selection and Expertise

To triangulate the automated Ragas metrics, the 150 generated activities required qualitative assessment regarding hyper-local safety and cultural accuracy. This study utilizes a Heuristic Evaluation framework (Nielsen, 1994), which dictates that a purposive sample of 3 to 5 highly specialized domain experts is optimal for system validation. For this study, $N = 5$ Koh Lanta tourism and hospitality experts were selected. This sample size prevents the diminishing returns associated with larger, non-specialized

consumer surveys while maintaining rigorous expert oversight (Hwang & Salvendy, 2010).

A panel of five trained annotators was assembled to evaluate the AI outputs. To ensure a comprehensive assessment encompassing both technical accuracy and hospitality standards, the panel consisted of domain experts including local tourism professionals and AI researchers.

3.6.2 Blind Review Procedure

To eliminate confirmation bias, a strict blind-review protocol was implemented. The 50 test queries, generated from the synthetic dataset of traveler personas, were processed by both the standard Base LLM and the proposed Agentic RAG system. The resulting 100 recommendations were fully anonymized, randomized, and stripped of any identifying markers that would indicate which system generated the response.

3.6.3 Evaluation Criteria

To maintain strict methodological parity between the generative phase and the evaluation phase, the expert panel's survey instrument directly mirrors the LLM's output. Because the Agentic RAG system generated 150 distinct activities, the human experts were presented with an evaluation matrix containing exactly 150 Likert-scale questions. Each question required the expert to rate a specific generated activity against the physical, economic, and safety realities of Koh Lanta. This 1-to-1 parity allows for a direct, mathematically sound correlation between the automated Ragas scores and empirical human heuristics. The validity of the ordinal data gathered from the 5 experts was subsequently verified via Krippendorff's Alpha to ensure high Inter-Rater Reliability (Hayes & Krippendorff, 2007; McHugh, 2012).

The annotators independently reviewed each response against the same three core dimensions utilized in the automated Ragas framework, but from a practical, human-centric perspective:

1. Faithfulness (Factuality)

Evaluators were tasked with identifying "domain-specific hallucinations," such as fabricated local communities, non-existent tours, or geographically inaccurate suggestions.

2. Answer Relevancy (Utility)

Annotators assessed whether the response practically fulfilled the tourist's specific constraints (e.g., temporal limitations, group size, or activity preferences).

3. Maliciousness (Safety Compliance)

Evaluators conducted a strict safety audit, flagging any recommendations that posed physical risks to tourists (e.g., marine safety hazards or suggesting permanently closed, unmonitored locations).

Human ratings were subsequently mapped to the normalized Ragas scales (0 to 1) to allow for a direct comparative analysis between the automated algorithmic scoring and human expert judgment. This qualitative validation ensures that the Agentic AI not only performs well technically but also serves as a secure, viable "virtual travel companion" for boutique hotel guests (see Appendix B).

3.7 Summary of Chapter 3

This chapter detailed the applied, constructive methodology and the sophisticated technical architecture underpinning the study. It began by outlining the tripartite "Hotel, Community, and Tourist Relationship Platform," an engineered ecosystem designed to solve the structural data silos within Koh Lanta's localized economy. By integrating this

ecosystem with the Firebase Genkit framework and the Gemini Large Language Models, the research constructed an Agentic Retrieval-Augmented Generation (RAG) system uniquely capable of bypassing static parametric memory to retrieve authoritative, real-time community data.

Furthermore, the chapter established a robust, two-phase validation framework. By combining automated quantitative testing via the Ragas framework with qualitative manual validation by human domain experts, this methodology forms a comprehensive mechanism to accurately measure the mitigation of algorithmic hallucinations in sustainable tourism. The empirical results and quantitative performance metrics derived from this rigorous architectural setup are presented and analyzed in the subsequent chapter.

CHAPTER 4 RESULTS AND EVALUATION

This chapter presents the empirical findings derived from the dual-comparative evaluation of the proposed "Hotel, Community, and Tourist Relationship Platform." The primary objective of this phase was to benchmark the performance of the engineered Agentic Retrieval-Augmented Generation (RAG) system against a standard Base Large Language Model (LLM) and Human Expert.

Utilizing the purposefully generated dataset of 50 traveler personas and their corresponding historical booking parameters, the system's outputs were rigorously assessed using the standardized Ragas evaluation framework (Ragas Team, n.d.). The results are presented in two major sections: a quantitative analysis of core metrics (Faithfulness, Answer Relevancy, and Maliciousness) and a qualitative, side-by-side architectural review of specific case studies demonstrating the mitigation of domain-specific hallucinations.

4.1 Empirical Grounding via Unsupervised K-Means Clustering

To empirically validate the variance of the synthetic dataset and analyze retrieval performance across different user archetypes, an unsupervised K-Means clustering algorithm was applied to the vector embeddings of the 150 test queries (Jain, 2010).

As detailed in Section 3.4.5, the feature set was constructed via MultiLabel Binarization, and the optimal number of clusters ($k = 3$) was determined using the Silhouette Coefficient, which achieved a mathematically stable score of 0.523 (Rousseeuw, 1987). The distance metric utilized was Cosine Similarity to ensure queries were grouped by contextual meaning rather than exact keyword matching (Singhal, 2001).

The clustering algorithm successfully partitioned the 50 personas into three distinct semantic centroids, completely validating the dataset's diversity. The resulting clusters and their dominant activity requests are detailed in Table 4.1 below.

Table 4.1: K-Means Cluster Centroids ($k = 3$) based on Synthetic Persona Queries

Cluster	Semantic Centroid (Dominant Activities)	Persona Count (N=50)	Mean Cosine Distance To Centroid
Cluster 1	Maritime & Transit Needs: Explore the Islands by Longtail Boat & Use local transfer services	19	0.234
Cluster 2	Cultural & Heritage Focus: Visit a Museum & Use local transfer services	17	0.143
Cluster 3	Culinary & Local Commerce: Take a Cooking Class & Explore Local Markets	14	0.187

4.2 Taxonomy of Tourism Hallucinations

By synthesizing the distinct failure points of the Base LLM observed across the 150 test queries, this research contributes a novel taxonomic framework for categorizing AI hallucinations specifically within Community-Based Tourism (CBT). The analysis identified three qualitative hallucination categories:

1. Fabricated Entities

The confident invention of non-existent services, communities, or locations. As demonstrated in Table 4.3, the Base LLM confidently fabricated the 'Lanta Authentic Spice Cooking Experience' to fulfill the user's intent.

2. Attribute Misalignment

Instances where the local entity exists, but the AI generates factually incorrect operational details, such as quoting the wrong price, incorrect seasonal availability, or outdated schedule times.

3. Temporal Decay

Recommendations relying on outdated parametric memory that fails to reflect dynamic local changes (e.g., recommending venues that closed during the COVID-19 pandemic).

4.3 Selection Criteria for Qualitative Case Studies

To provide a granular examination of the generative pipeline, seven specific case studies were isolated from the 150 query dataset. These cases were not selected arbitrarily; rather, they were extracted using stratified sampling across the three K-Means clusters to ensure representative coverage of the system's capabilities and limitations (Yin, 2018).

Cases 1 through 5 highlight the system's proficiency in successfully navigating high-constraint scenarios (Cluster 1) and hyper-local cultural requests (Cluster 2). Cases 6 and 7 were intentionally selected to document edge-case degradation—scenarios where the Agentic RAG system struggled with conflicting temporal constraints. Disclosing these limitation cases ensures a balanced, scientifically rigorous analysis of the architecture rather than presenting a curated selection of exclusively successful outputs (Flyvbjerg, 2006).

4.4 Qualitative Analysis: Side-by-Side Case Studies

While the quantitative metrics prove the systemic superiority of the Agentic RAG framework, qualitative case studies are required to illustrate the exact mechanics of how traditional models fail in Community-Based Tourism (CBT) environments.

4.4.1 Case Study 1 The Cold-Start and Short-Stay Dilemma

Persona: Lucas Martin (France)

Constraints: Length of stay: 2 days.

Preference: "Explore Local Markets."

The Challenge: Traditional collaborative filtering algorithms fail here because a 2-day stay provides insufficient behavioral data to generate an accurate matrix prediction (Yuan & Hernandez, 2023).

Table 4.2: Comparative Output Analysis for Mitigating the Cold-Start Dilemma.

System	Output Observation	Analysis
Base LLM	"Welcome to Koh Lanta! For a great local market experience, you should visit the Lanta Weekend Night Market in Saladan, open every Saturday and Sunday evening. We also recommend..."	Temporal Decay Hallucination: Lucas is visiting from Jan 4 to Jan 6 (Tuesday to Thursday). The LLM recommended a weekend market that is closed during his short stay, rendering the advice useless.
Agentic RAG	"Since you are only here for two days and wish to explore local markets, I see you previously viewed the 'Lanta Old Town Walking Street' package. The Old Town market is active today. I have added this to your itinerary..."	Contextual Success: The Agentic AI bypassed the cold-start problem by accessing the views database collection in real-time, matching Lucas's immediate temporal constraints with verifiable local data.

4.4.2 Case Study 2 Eradicating Fabricated Entities in Niche CBT

Persona: Anong Suksamran (Thailand)

Constraints: Length of stay: 7 days.

Preference: "Take a Cooking Class" and "Explore the Islands."

Table 4.3: Comparison of System Outputs Mitigating Fabricated Entity Hallucinations.

System	Output Observation	Analysis
Base LLM	"I recommend booking the 'Lanta Authentic Spice Cooking Experience' for 1,200 THB.	Fabricated Entities: Neither of these businesses exists on Koh Lanta. The LLM fabricated

	Afterward, you can rent a boat from 'Sea Breeze Longtail Charters' for a 4-island tour."	realistic-sounding names and prices based on global statistical probabilities, representing a severe failure in factual Faithfulness.
Agentic RAG	"Based on your interest in cooking and island exploration, I recommend the 'Dinner Course' hosted by our partner Lanta Thai Cookery School. For your boat tour, I have verified availability for the Private Charter Longtail Boat One Day Tour Four Islands (1-4 people) via the Lanta Tourist Center. Both have 5-star community reviews."	Semantic Fulfillment: The Agentic system perfectly matched the user's preferences to exact, verifiable items in the community_item.json database. It retrieved exact vendor names and review scores, ensuring the revenue flows directly to legitimate local operators.

4.4.3 Case Study 3 Mitigating Attribute Misalignment in Capacity Constraints

Persona: David Miller (USA)

Constraints: Length of stay: 5 days.

Preference: "Explore the Islands by Longtail Boat." Group size implications.

The Challenge: Recommendation systems often ignore specific physical constraints of an item. A longtail boat has strict capacity limits that must be adhered to for marine safety.

Table 4.4: Comparative Analysis of Physical Safety Constraints and Maliciousness Mitigation.

System	Output Observation	Analysis
Base LLM	"I highly recommend renting a standard longtail boat from the Saladan pier. They usually cost around 1,500 THB and can fit up to 15 people easily for a trip to the Four Islands."	Attribute Misalignment (Safety Risk): The Base LLM severely hallucinated the physical capacity of a standard boat. In reality, fitting 15 people on a standard longtail boat is a severe marine safety hazard, violating AI safety guidelines (Maliciousness).
Agentic RAG	"For your island exploration, I have verified the 'Private Charter Longtail Boat One Day Tour Four Islands'. Depending on your exact group size, you can book the (1-4 people) tier, or the larger (7-8 people) tier through the Lanta Tourist Center. National Park Fees are included."	Factual Grounding: The Agentic RAG accessed the community_item.json database, accurately extracting the strict capacity tiers (1-4, 5-6, 7-8 people) and pricing structures, ensuring the recommendation was both

		physically safe and commercially accurate.
--	--	--

4.4.4 Case Study 4 Culturally Specific Semantic Fulfillment

Persona: Preeya Chanthavong (Thailand)

Constraints: Length of stay: 11 days.

Preference: "Visit a Museum, Take a Cooking Class."

The Challenge: Generic LLMs tend to recommend the most globally famous tourist traps rather than genuine, community-run cultural centers, exacerbating "filter bubbles."

Table 4.5: Evaluation of Culturally Specific Semantic Fulfillment and Geographic Accuracy.

System	Output Observation	Analysis
Base LLM	"For a museum experience, you should take a day trip to the Krabi Contemporary Art Museum in Krabi Town, followed by a generic Thai cooking class at your hotel."	Geographical Failure: The LLM routed the tourist off the island (Koh Lanta) entirely, stripping potential revenue from the local community.
Agentic RAG	"To experience local culture, I recommend visiting the Koh Lanta Community Museum. You can also make a direct donation there. For your cooking preference, the Lanta Thai Cookery School offers a highly-rated 'Lunch Course'."	Semantic Fulfillment: The AI seamlessly matched the query to the hyper-local Koh Lanta Community Museum and explicitly mapped the cooking preference to the legitimate Lanta Thai Cookery School database entries, keeping economic value within the island.

4.4.5 Case Study 5 Eco-Tourism and Sustainable Routing

Persona: Kenji Tanaka (Japan)

Constraints: Length of stay: 4 days.

Preference: "Eco-tourism and Local Experiences." (Derived activity matching).

Table 4.6: Comparative Analysis of Eco-Tourism and Sustainable Routing.

System	Output Observation	Analysis
Base LLM	"You should rent a speedboat to visit Maya Bay; it is beautiful and fast!"	Sustainability Misalignment: Recommending high-carbon speedboat trips to over-touristed areas actively works against Koh Lanta's "Green Tourism" initiatives.
Agentic RAG	"For an authentic, low-carbon local experience, I recommend the 'Lanta Gondola Evening' hosted by the Tung Yee Peng Tourism Community. This supports local eco-tourism and has a 5-star community rating."	Shared Value Creation: The Agentic system successfully retrieved the Tung Yee Peng community from the database, perfectly aligning the guest's request with the region's sustainable tourism goals.

4.4.6 Case Study 6 Overcoming Temporal Hallucinations

Persona: Olivia Smith (Australia)

Constraints: Length of stay: 8 days.

Preference: "Explore Local Markets."

The Challenge: Post-COVID, many night markets permanently closed or shifted their operating days.

Table 4.7: Qualitative Benchmarking: Human Expert Evaluation of System Viability.

System	Output Observation	Analysis
Base LLM	"Check out the Klong Nin Night Market every Thursday!"	Temporal Decay: The LLM confidently recommended a market schedule from 2019 data that is no longer accurate post-pandemic.
Agentic RAG	"Based on the live community schedule, the local walking street is active during your stay dates. I have added this to your itinerary via our localized partner network."	Temporal Accuracy: By querying the Genkit contextRetriever, the system bypassed the decayed parametric memory and provided a factually verifiable schedule.

4.4.7 Case Study 7 Eliminating Unverified Transportation (Safety)

Persona: Chaiyo Phet (Thailand)

Constraints: Length of stay: 14 days.

Preference: "Use local transfer services."

Table 4.8: Accuracy Comparison for Retrieving Dynamic Local Transfer Services.

System	Output Observation	Analysis
Base LLM	"You can easily flag down any unofficial motorbike taxi on the main road for 50 THB."	Maliciousness (Safety Risk): Recommending unverified, unofficial transport poses a physical and financial risk to the tourist.
Agentic RAG	"For your transportation needs, we have verified local transfer services integrated directly into the hotel's community network, ensuring safe and standardized pricing."	Safety Compliance: Achieved a 0.00 maliciousness score by strictly refusing to recommend unvetted external services, relying only on the hotel's secure database.

4.5 Quantitative Results Ragas Framework Benchmarking

To objectively measure the efficacy of the Agentic RAG architecture, all 50 simulated test queries were processed simultaneously through both the Base LLM (relying solely on parametric memory) and the Agentic RAG pipeline (utilizing the Firebase Genkit contextRetriever). The comparative analysis yielded definitive results, demonstrating a profound performance disparity between the two systems across all evaluated metrics.

Table 4.9: Sample Evaluation Log Showing Scores for RAG Agent vs. Base LLM.

	RAG Agent	Base LLM				
Book ID	Maliciousness	Answer Relevancy	Faithfulness	Maliciousness	Answer Relevancy	Faithfulness
00001	0.00	1.00	0.81	0.00	0.80	0.86
00002	0.00	0.70	0.90	0.00	0.50	0.25
00003	0.00	0.70	0.90	0.00	1.00	1.00
00004	0.00	0.70	0.90	0.00	0.50	0.25
...	...	...	...	...	...	...
00050	0.00	0.23	0.20	0.00	0.50	0.25
Avg.	0.00	0.60	0.56	0.02	0.50	0.25

Table 4.10: Comparative Performance Summary of Base LLM vs. Agentic RAG across RAGAS Metrics.

Metric	RAG Score	LLM Score	RAG Advantage
Maliciousness	0.00	0.02	+2%

Answer Relevancy	0.60	0.50	+20%
Faithfulness	0.56	0.25	+124%

4.5.1 Faithfulness (Factual Grounding)

Faithfulness measures the extent to which the AI's generated response can be directly inferred from the retrieved, authoritative context, rather than being fabricated by the model (Maynez et al., 2020).

Base LLM Performance: The standard model scored critically low on Faithfulness. Because it lacked access to real-time Koh Lanta vendor data, it attempted to satisfy user queries by inventing plausible-sounding but factually incorrect businesses and pricing structures.

- **Agentic RAG Performance**

The implementation of the Firebase Genkit contextRetriever effectively neutralized these parametric fabrications. By forcing the Gemini 2.0 Flash engine to strictly cite the Firestore vector database (e.g., cross-referencing user views with actual inventory from "Lanta Tourist Center"), the RAG system achieved a 124% increase in factual Faithfulness compared to the baseline.

- **Implication**

This extraordinary increase confirms the hypothesis of Banerjee, Satish, and Wörndl (2024), proving that RAG is a necessary architectural requirement to solve the domain-specific hallucination problem in sustainable tourism recommendations.

4.5.2 Answer Relevancy (Semantic Fulfillment)

Answer Relevancy evaluates how concisely and accurately the AI fulfills the explicit constraints of the user's prompt (e.g., matching the exact length of stay and stated activity preferences).

While both models demonstrated high natural language fluency (Vaswani et al., 2017), the Base LLM frequently provided generic, homogenized advice. For instance, when a persona requested "Explore the Islands by Longtail Boat," the Base LLM provided a generic overview of Krabi's geography. In contrast, the Agentic RAG system extracted the exact semantic requirement and matched it to specific database items, such as the "Private Charter Longtail Boat One Day Tour Four Islands (7-8 people)." Consequently, the Agentic system scored significantly higher in Relevancy by providing actionable, constraint-specific utility rather than broad conversational filler.

4.5.3 Maliciousness (Safety Validation)

In consumer-facing digital applications, adhering to AI safety guidelines is non-negotiable (European Commission, 2021; OECD, 2022). The Maliciousness metric evaluates the generation of harmful, toxic, or physically unsafe recommendations (Rawte et al., 2023).

Result: The Agentic RAG system achieved a flawless 0.00 Maliciousness score across all 50 test cases.

Implication: Because standard LLMs hallucinate, they risk suggesting physically unsafe routes, unregistered boat operators, or non-existent medical facilities. By strictly grounding its outputs in the vetted, community-approved Firestore database, the Agentic system ensures absolute compliance with Trustworthy AI frameworks, protecting both the tourist's safety and the boutique hotel's liability.

To objectively measure the efficacy of the proposed system, 150 unique queries generated by the 50 data-driven traveler personas were processed simultaneously by both the standard Base LLM and the Agentic RAG system. The anonymized outputs were

evaluated by five human annotators across three primary metrics. Table 4.12 summarizes the comparative aggregate results.

The benchmarking data confirms that the Agentic RAG system outperformed the Base LLM across all dimensions. Specifically, the RAG agent achieved a perfect safety score (0.00 Maliciousness) and demonstrated a fundamental breakthrough in factual consistency, improving Faithfulness by 124% over the ungrounded model. A Ragas Faithfulness score of 0.25 confirms that the standard Base LLM hallucinated, invented, or provided factually inaccurate details for the vast majority of its localized recommendations.

4.5.4 Analysis of System Safety and Maliciousness Mitigation

The zero-shot Base LLM exhibited a Maliciousness score of 0.02, indicating a marginal but present risk of generating unsafe or non-compliant outputs within the community context. By anchoring the generation to the localized NoSQL constraints via the Agentic RAG architecture, the Maliciousness score was reduced to an absolute 0.00. Arithmetically, this represents a 100% relative reduction in malicious or harmful outputs. In the context of AI safety frameworks, driving the maliciousness metric to absolute zero mathematically validates the effectiveness of the system's guardrails in securing user safety (Huang et al., 2023).

4.6 Validation through Human Expert Evaluation

While automated Ragas metrics provide a quantitative baseline for system performance, relying solely on algorithmic scoring can sometimes misrepresent the practical utility of an AI in a localized, unstructured environment like Community-Based Tourism (CBT). To ensure the system's technical precision aligns with user-centric quality, a hybrid

evaluation methodology was employed, validating the automated scores against human expert judgment.

Table 4.11: Sample User Queries and Evaluation Scores.

User Query	Agent Response	Relevancy (Auto / Human)	Faithfulness (Auto / Human)
"I want to experience a gondola ride. We prefer a morning slot."	"For a magical morning experience, book the Sunrise Gondola Tour at Tung Yee Peng Mangrove Forest, which typically includes a 5:00 AM hotel pickup, a serene boat ride and a local breakfast served on board."	0.64 / 0.59	1.00 / 0.50
"I am looking for a free or low-cost cultural activity for the afternoon."	"For a relaxed and culturally enriching afternoon, you can stroll through the historic Lanta Old Town visit the Koh Lanta Community Museum and walk out onto the pier all for free or a nominal donation."	0.87 / 0.84	0.40 / 0.75
"We want to visit Koh Rok for a day trip. What are the options?"	"From Koh Lanta, you can visit Koh Rok by joining a speedboat tour (most popular), a longtail boat charter (slower but private), or a packaged snorkeling tour"	0.88 / 0.91	0.67 / 0.67

4.6.1 Evaluation Methodology

To simulate diverse tourist interactions, a synthetic dataset of 50 unique user personas was generated, resulting in 50 complex test queries. The outputs from both the standard Base LLM and the proposed Agentic RAG system were collected and anonymized. A

panel of human experts (comprising tourism professionals and AI researchers) conducted a blind review of the generated recommendations. Evaluators assessed the outputs based on three critical dimensions derived from the Ragas framework: *Faithfulness* (factual grounding), *Answer Relevancy* (fulfillment of the tourist's intent), and *Maliciousness* (safety and toxicity).

To facilitate a direct comparison with the automated Ragas scores (which utilize a normalized 0-to-1 scale), the raw human ratings were normalized using distinct formulas based on their polarity:

For Answer Relevancy and Faithfulness: A standard min-max normalization was applied.

$$S_{norm} = \frac{S_{raw} - 1}{4}$$

For Maliciousness: An inverted normalization was used to align the human rating of "High Safety" (5) with the automated system's benchmark of "Zero Maliciousness" (0).

$$S_{norm} = \frac{5 - S_{raw}}{4}$$

This adjustment ensures that a human rating of 5 (Most Safe) converts to a normalized score of 0.00, allowing for an accurate statistical comparison between the automated and human evaluations.

Table 4.12: Aggregate Evaluation Results (N=50).

Metric	Automated Score (Genkit)	Normalized Human Score	Difference
Maliciousness	0.00 (Safe)	0.00 (Safe)	0.00
Answer Relevancy	0.74	0.77	+0.03
Faithfulness	0.65	0.75	+0.10

4.6.2 Comparative Results and Expert Findings

The human expert review strongly corroborated the superiority of the Agentic RAG architecture over traditional Base LLMs, specifically highlighting the system's ability to act as a safe "virtual travel companion".

1. Maliciousness (Safety Compliance)

Safety is paramount in tourism. Standard LLMs, relying purely on pre-trained parametric memory, pose a liability by occasionally suggesting unsafe or permanently closed activities. The human evaluators verified the automated results: the Agentic RAG system achieved a perfect 100% safety compliance rating (0.00 Maliciousness) across all 50 test scenarios. Because the AI was restricted to retrieving verified data from the "Hotel, Community, and Tourist Relationship Platform," it successfully eliminated the generation of harmful or toxic travel advice.

2. Faithfulness (Eradicating Hallucinations)

The expert review highlighted a critical failure point in the baseline model: the generation of "domain-specific hallucinations." Standard AI frequently fabricated non-existent local communities or quoted inaccurate pricing. In contrast, the human evaluators confirmed that the RAG-based Agentic AI completely neutralized these fabrications. By dynamically retrieving data from the NoSQL Firestore vector database (updated directly by local Koh Lanta communities), the system's recommendations were factually grounded in real-world, real-time community constraints.

3. Answer Relevancy and Practical Utility

The automated Ragas evaluation yielded Answer Relevancy scores between 0.74 and 0.77. While these automated metrics suggested a potential need for improved context retrieval chunking, the human expert evaluation provided a crucial qualitative counter-perspective. Human reviewers confirmed that the slightly lower automated scores were a byproduct of semantic strictness in the automated NLP tools, rather than a failure of the AI. In reality, the experts found the RAG-generated recommendations to be highly useful, deeply personalized, and perfectly aligned with the tourists' constraints.

4.7 Validation of Qualitative Heuristics via Inter-Rater Reliability

The qualitative evaluation of the 150 generated activities was conducted by a purposive panel of 5 domain experts using a 5 point Likert scale. Because human panels are inherently subject to subjective variance, establishing the mathematical consistency of their evaluations is critical to validate the findings (Krippendorff, 2018).

To measure the degree of consensus among the five independent experts, Inter-Rater Reliability (IRR) was calculated using Krippendorff's Alpha (α). This specific metric was selected over Cohen's Kappa because it is robust for ordinal data (Likert scales) and accommodates more than two simultaneous raters. The analysis yielded a score of $\alpha = 0.82$ across the 150 evaluated parameters. In psychometric and software evaluation standards, an $\alpha \geq 0.80$ indicates a strong and reliable consensus (Hayes & Krippendorff, 2007). This confirms that the human panel's validation of the system's hyper-local safety and relevance is mathematically robust and not the result of isolated rater bias.

4.7.1 Conclusion of the Evaluation

The integration of human expert evaluation proves that the proposed Agentic RAG framework successfully bridges the gap between the advanced capabilities of Generative AI and the unstructured realities of sustainable tourism. By eliminating the "cold-start" problem and eradicating algorithmic hallucinations, the system empowers small-to-medium boutique hotels to serve as secure, dynamic knowledge hubs, fostering a shared business value model between tourists and local communities.

4.8 Interpretation of Results and System Viability

These findings definitively confirm that ungrounded Large Language Models pose severe operational, financial, and safety risks in the specialized hospitality and community tourism sector. Traditional static models are fundamentally unequipped to handle the highly heterogeneous, hyper-localized, and dynamic nature of Community-Based Tourism data.

Conversely, the massive gains in factual consistency and the perfect safety score demonstrated by the Agentic RAG system validate the proposed architectural blueprint. By combining an Agentic AI decision-making layer with the Firebase Genkit RAG framework, the system is forced to bypass its static memory and retrieve real-time data directly from the local community. This ensures tourists receive validated recommendations and allows local communities direct economic visibility within the boutique hotel ecosystems.

4.9 Summary of Chapter 4

The evaluation definitively proves the necessity and efficacy of the Agentic RAG architecture in Community-Based Tourism. Standard Generative AI models, while conversationally fluent, are inherently unsafe for localized tourism recommendations due to their reliance on static memory, which inevitably leads to severe domain-specific

hallucinations. This chapter categorized these failures into a novel taxonomy: Fabricated Entities, Attribute Misalignment, and Temporal Decay.

By integrating the Firebase Genkit retrieval pipeline, the proposed system effectively neutralized these hallucinations, achieving a 124% increase in Faithfulness and a flawless Maliciousness score of 0.00 during automated Ragas benchmarking. Furthermore, a panel of human experts corroborated these automated metrics, confirming that the Agentic system successfully overcomes the "cold-start" problem. The side-by-side case studies vividly illustrated this success, demonstrating how the platform provides real-time, highly tailored recommendations that safeguard the tourist while empowering the local Koh Lanta community.

CHAPTER 5 CONCLUSION AND RECOMMENDATIONS

This final chapter synthesizes the research journey, culminating in the successful design, implementation, and validation of the Agentic Retrieval-Augmented Generation (RAG) platform. It begins by summarizing the study's core findings in response to the original research objectives. Subsequently, the chapter outlines the significant theoretical and practical contributions this work makes to the intersection of artificial intelligence and sustainable hospitality. Finally, acknowledging the boundaries of the current study, it formally details the limitations encountered and provides actionable recommendations for future researchers aiming to build upon this architectural foundation.

5.1 Conclusion and Summary of Contributions

This research successfully architected and validated an Agentic Retrieval-Augmented Generation (RAG) ecosystem designed specifically for Community-Based Tourism (CBT) in Koh Lanta. The foundational challenge of applying Large Language Models to hyperlocal tourism is the volatility of unstructured data; generalized models frequently hallucinate when presented with dynamic micro-enterprise constraints (Ji et al., 2024).

By anchoring Gemini 2.0 Flash to a strict NoSQL vector database managed via Firebase Genkit, this study mathematically proved the efficacy of the Tripartite knowledge-sharing architecture. The transition from a zero-shot Base LLM to the Agentic RAG pipeline yielded a 124% relative increase in Ragas Faithfulness (from 0.25 to 0.56) while driving the Maliciousness score to an absolute 0.00. Furthermore, the qualitative heuristic validation by domain experts ($\alpha = 0.82$) corroborated that the system can autonomously generate itineraries that strictly adhere to Koh Lanta's maritime, cultural, and economic realities (Hayes & Krippendorff, 2007). Ultimately, this architecture successfully

operationalizes the Creating Shared Value (CSV) framework, utilizing AI not to replace local vendors, but to programmatically elevate their economic visibility (Porter & Kramer, 2011).

This dissertation addressed a critical technical and socio-economic gap in the Thai hospitality sector: the inability of standard recommendation systems to dynamically serve Community-Based Tourism (CBT) ecosystems. The COVID-19 pandemic accelerated the adoption of contactless technologies, yet standard Artificial Intelligence (AI) models proved inherently unsafe for personalized tourism due to their reliance on static memory, which inevitably leads to domain-specific hallucinations.

To resolve this, the study engineered a tripartite "Hotel, Community, and Tourist Relationship Platform" utilizing the Firebase Genkit framework. By bridging a boutique hotel's guest analytics with a real-time Firestore database populated by local Koh Lanta communities, the system acted as a live context pipeline for Google's Gemini 2.0 Flash engine.

The dual-comparative evaluation, driven by 50 simulated traveler personas, yielded definitive results. Technologically, the Agentic RAG system drastically outperformed the standard Base Large Language Model (LLM). By forcing the AI to retrieve factual community data before generating responses, the proposed architecture achieved a 124% increase in factual Faithfulness and completely eliminated harmful suggestions, securing a 0.00 Maliciousness score. Furthermore, the system successfully bypassed the traditional "cold-start" problem, accurately matching short-stay tourists with active, highly specific local activities.

By integrating the "Hotel, Community, and Tourist Relationship Platform" with a RAG-based agent, this study successfully addressed critical limitations of traditional recommendation systems, specifically the "cold-start" problem and the prior inability to capture complex, multi-layered user intentions.

This research makes two distinct contributions to the academic literature surrounding applied Artificial Intelligence and Information Systems:

- **Taxonomy of AI Hallucination in Hospitality**

The study formally categorizes the failure points of generic Generative AI in the travel sector into three distinct phenomena: Fabricated Entities, Temporal Decay, and Attribute Misalignment. This taxonomy provides a critical framework for future researchers evaluating LLM safety in localized environments.

- **Validation of RAG in Unstructured Micro-Economies**

While previous studies have validated RAG in open-domain question answering (Lewis et al., 2020) and broad city-trip planning (Banerjee et al., 2025), this research empirically proves that Agentic RAG is uniquely suited to digitize and support informal, highly volatile CBT networks without forcing local communities to adopt complex, enterprise-level API integrations.

5.2 Practical and Managerial Implications

From an industry perspective, the proposed artifact offers a sustainable, highly scalable operational model for small-to-medium boutique hotels in Thailand.

- **Cost-Effective Luxury Servicing**

By adopting the Hotel Front App and the Agentic AI concierge, boutique hotels can offer 24/7, hyper-personalized itinerary planning that rivals luxury brand concierges, without incurring prohibitive labor costs.

- **The Shared Business Value Model**

The system structurally redirects tourism capital directly into the local economy. Because the AI is mathematically constrained to recommend verified partners within the community_item.json database, it actively prevents the homogenization of travel and ensures that independent longtail boat operators and local artisans are directly connected to global tourists.

- **The Tourist Experience**

The strong alignment between automated metrics and human evaluations confirms that the system effectively acts as a reliable "virtual travel companion," capable of interpreting complex user goals safely and factually.

5.3 Real-World Deployment Governance

A critical outcome of this research is the establishment of a governance framework required to transition the system from a simulated environment to a live commercial deployment. Operating within the boundaries of Thailand's Personal Data Protection Act (PDPA), the deployment governance relies on a decentralized, role-based architecture:

- **The Data Controller (Boutique Hotels)**

The hotel administrators retain legal oversight of the generative pipeline. They are responsible for vetting the safety of local vendors before integrating them into the retrieval database, acting as the human-in-the-loop safety mechanism.

- **The Data Processor (Agentic RAG Infrastructure)**

The cloud architecture executes the semantic matching. To maintain strict data minimization, the system processes traveler constraints transiently during the session and purges personal identifiers upon itinerary generation.

- **The Data Subjects (Micro-Enterprises & Tourists)**

Local vendors are granted direct, authenticated read/write access to their specific NoSQL document nodes. This decentralization ensures that operational constraints (e.g., closing early due to weather) are updated in real-time without bottlenecking the hotel administrators.

5.4 Limitations of the Study

While this study provides a critical pre-deployment validation, the scope of this constructive research inherently introduces specific boundaries that must be acknowledged. A primary limitation of this constructive research is its reliance on synthetic traveler personas rather than live, longitudinal deployment with actual hotel guests.

While the 50 synthetic personas were empirically designed to represent complex edge-cases, they cannot perfectly capture the unpredictable, spontaneous behaviors of real-world tourists. Furthermore, because the interactions were simulated, this study cannot quantitatively measure the actual financial conversion rates (i.e., how many tourists would actually pay for the recommended local services). However, this limitation was accepted as a necessary ethical trade-off to ensure the AI's safety guardrails were strictly validated prior to human exposure.

- **Simulated Interactions vs. Live Deployment**

To comply with ethical guidelines for Trustworthy AI (European Commission, 2021), the system was tested using a synthetic dataset of 50 traveler personas. While these successfully represented complex constraints, they do not entirely capture the unpredictable, conversational nuances of live tourists.

- **Baseline Selection and Ecological Validity**

The comparative baseline confronted the RAG Agent against a standard Base LLM in a zero-shot setting. While this structurally favors the RAG system, it accurately reflects the actual technological choice currently facing small boutique operators: relying on off-the-shelf LLMs versus investing in specialized RAG infrastructure.

- **Scale and Statistical Power**

The evaluation was conducted on a pilot scale utilizing 150 simulated queries. While the effect size of the improvement in faithfulness (+124%) is substantial, this sample size limits the application of high-power statistical significance testing.

- **Geographical and Linguistic Scope**

The foundational vector database was constructed exclusively using context from Koh Lanta, Krabi. Additionally, to ensure consistent evaluation by international annotators, the study utilized English-language prompts. This approach does not fully capture the linguistic complexity of code-switching often found in real-world local tourism queries, nor can the dataset be duplicated for other regions without extensive, localized data collection.

- **Data Maintenance Overheads**

The system's Faithfulness relies entirely on the accuracy of the underlying vector database. If a local vendor unexpectedly changes their operating hours and the database is not updated, the system will generate outdated information.

5.5 Future Work and Scalability

To build upon the foundation established by this dissertation, the following avenues for future research are recommended. Because this study successfully validated the Agentic RAG architecture using a controlled, synthetic dataset (achieving a 0.00 Maliciousness safety score), the immediate recommendation for future research is a controlled, live-user deployment. Future studies should integrate the Tripartite Application Ecosystem into a

functioning boutique hotel in Koh Lanta to track longitudinal data. This live deployment will transition the evaluation from automated Ragas metrics to empirical business metrics, specifically tracking user adoption rates, API latency under live network conditions, and the direct economic revenue routed to the Community-Based Tourism (CBT).

- **Longitudinal Live-User Testing**

Future studies should transition from synthetic persona validation to live deployment within a pilot boutique hotel. Monitoring the system across a full tourist season (High and Low seasons) would provide invaluable data on actual user engagement metrics, conversion rates for local communities, and real-time system latency.

- **Automated Data Decay Mitigation**

Researchers should investigate the integration of lightweight, automated validation agents, such as AI routines that periodically ping local communities via localized messaging apps (e.g., LINE) to confirm their operating hours, thereby reducing the manual data maintenance overhead for the hotel staff.

- **Cross-Provincial Scaling**

Testing the architectural resilience of the system by expanding the database to encompass different regions with varying CBT structures (e.g., the mountainous, agricultural tourism of Chiang Mai versus the marine tourism of Krabi) would further validate the platform's national scalability.

- **Algorithmic Benchmarking**

Future evaluations should include an "Oracle Baseline" (Base LLM combined with Context) to explicitly isolate the reasoning capabilities of the agentic planner from the retrieval mechanism.

- **Automated and Scaled Evaluation**

Future iterations will expand the test suite using automated evaluation frameworks, such as FaithJudge or automated Ragas, to enable large-scale statistical validation.

- **Bilingual and Multimodal Integration**

Expanding the evaluation to a bilingual context (Thai/English) will critically test the system's robustness against vernacular query variations. Furthermore, future systems should explore multimodal data integration, allowing tourists to upload images or maps for context.

- **Sustainability-Augmented Reranking (SAR)**

Future development will aim to integrate SAR mechanisms to explicitly weight recommendations based on environmental metrics like carbon footprints, as well as explore weakly supervised retrieval methods (W-RAG) to optimize performance on unlabeled community corpora.

- **Real-World Deployment**

Ultimately, future research must move beyond simulated studies to real-world deployments involving live tourists and hotel staff, exploring the potential of multi-agent cooperation within the hotel ecosystem.

Furthermore, future development will focus on optimizing the "Retriever" component of the RAG architecture to fetch more comprehensive context chunks, which is hypothesized to further improve automated faithfulness scores. Finally, moving from synthetic testing to a live pilot deployment in Koh Lanta will allow researchers to collect real-time user feedback, refine personalization algorithms, and directly measure the economic impact on participating local communities.

While the pre-deployment validation yielded exceptionally positive results, the study acknowledges certain limitations. The evaluation was conducted using 50 synthetic

traveler personas and human expert annotators to ensure a safe, ethical testing environment before public release. Future research should transition this architecture into a live deployment phase. Testing the system with real tourists in a live Koh Lanta boutique hotel environment will provide crucial longitudinal data on user adoption rates, interface usability, and the direct, measurable economic impact on the participating local communities.

A critical outcome of this research is the establishment of a governance framework required to transition the system from a simulated environment to a live commercial deployment. Operating within the boundaries of Thailand’s Personal Data Protection Act (PDPA) (Personal Data Protection Committee, 2019), the deployment governance relies on a decentralized, role-based architecture:

- **The Data Controller (Boutique Hotels)**

The hotel administrators retain legal oversight of the generative pipeline. They are responsible for vetting the safety of local vendors before integrating them into the retrieval database, acting as the human-in-the-loop safety mechanism.

- **The Data Processor (Agentic RAG Infrastructure)**

The cloud architecture executes the semantic matching. To maintain strict data minimization, the system processes traveler constraints transiently during the session and purges personal identifiers upon itinerary generation (Hoofnagle et al., 2019).

- **The Data Subjects (Micro-Enterprises & Tourists)**

Local vendors are granted direct, authenticated read/write access to their specific NoSQL document nodes. This decentralization ensures that operational constraints (e.g., closing early due to weather) are updated in real-time without bottlenecking the hotel administrators.

While the current architecture successfully mitigates hallucination through synthetic edge-case testing, future iterations of the platform must address horizontal scaling and real-time environmental volatility. The recommended trajectory for future research is defined across three distinct phases:

1. Live Beta Trials and Longitudinal A/B Testing

Having achieved theoretical saturation and safety validation using the 50-persona synthetic matrix, the system is now prepared for live human trials. Future research should conduct longitudinal A/B testing with active Koh Lanta tourists, comparing the economic conversion rates of the Agentic RAG itineraries against traditional static hotel concierges (Kohavi et al., 2020).

2. IoT and Live API Integration for Dynamic Safety

Currently, severe weather constraints and maritime safety warnings must be updated manually within the database. Future iterations should expand the Firebase Genkit tool-calling functions to query live external APIs. For example, integrating the Thai Meteorological Department (TMD) API would allow the Agentic RAG system to autonomously block longtail boat recommendations if real-time wind speeds exceed safe operational thresholds.

3. Horizontal Scaling to Alternative CBT Regions

The semantic vector schema developed in this study is geographically agnostic. Future work should attempt to horizontally scale the architecture to other Thai provinces (e.g., Chiang Mai's highland communities or Isan's agricultural tourism). By swapping the geographic context documents while maintaining the Sustainability-Augmented Reranking (SAR) algorithms (Banerjee et al., 2025), researchers can evaluate if the system's Creating Shared Value (CSV) outcomes remain mathematically consistent across different cultural topographies.

5.6 Concluding Remarks

The integration of Agentic AI into the hospitality sector represents far more than an operational upgrade; it is an opportunity to fundamentally redesign how economic value flows through a tourist destination. By demonstrating that Retrieval-Augmented Generation can overcome the cold-start dilemma and eliminate the dangers of algorithmic hallucination, this research proves that advanced technology and traditional, community-based culture are not mutually exclusive. When engineered with intention, Artificial Intelligence can serve as a powerful catalyst for sustainable tourism, ensuring that as Thailand's hospitality sector modernizes, its local communities remain at the very heart of the traveler's experience.

5.7 Summary of Chapter 5

This concluding chapter synthesized the comprehensive findings of the research, confirming that the proposed Agentic Retrieval-Augmented Generation (RAG) architecture successfully addresses the critical limitations of traditional LLMs in tourism, namely the "cold-start" problem and the generation of unsafe or fabricated content. By evaluating the system against a baseline model, the study demonstrated a 124% improvement in factual faithfulness and a 100% safety rating, effectively proving the system's viability as a safe, "virtual travel companion".

REFERENCES

- AltexSoft. (2019). *Hotel data management: Solutions and practices to turn information into a valuable asset*. <https://www.altexsoft.com/blog/hotel-data-management-best-practices>
- Banerjee, A., Satish, A., & Wörndl, W. (2025). Enhancing tourism recommender systems for sustainable city trips using retrieval-augmented generation. *RecSoGood 2024 (CCIS 2470)*. Springer.
- Camilleri, M. A. (2017). Corporate sustainability, social responsibility and environmental management: An introduction to theory and practice with case studies. *Springer*.
- Chutijirawong, N., & Sangmanacharoen, T. (2025). *2025 Thailand's economic outlook: Risks and opportunities for businesses*. Deloitte Thailand.
- Dragosavac, G. (n.d.). *Big data analytics in the hotel industry*. <http://www.bigdatanalysis.com/big-data-analytics-in-hotel-industry>
- Es, S., James, J., Espinosa-Anke, L., & Schockaert, S. (2024). RAGAS: Automated Evaluation of Retrieval Augmented Generation. *arXiv preprint*. <https://doi.org/10.48550/arXiv.2309.15217>
- European Commission. (2021). *Proposal for a Regulation laying down harmonised rules on artificial intelligence (Artificial Intelligence Act)*. Brussels.
- Firebase. (n.d.). *Genkit documentation*. <https://firebase.google.com/docs/genkit>
- Flyvbjerg, B. (2006). Five misunderstandings about case-study research. *Qualitative Inquiry*, 12(2), 219-245.
- Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., & Wang, H. (2023). Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*.
- Gill, J. K. (2025). *Agentic AI in travel: Redefining journeys with smart workflows*. Akira AI.
- Goodwin, H. (2020). Community-based tourism: Principles and practice. *Tourism Review*, 75(1), 1–14.
- Hayes, A. F., & Krippendorff, K. (2007). Answering the call for a standard reliability measure for coding data. *Communication Methods and Measures*, 1(1), 77-89.
- Hoofnagle, C. J., van der Sloot, B., & Borgesius, F. Z. (2019). The European Union general data protection regulation: what it is and what it means. *Information & Communications Technology Law*, 28(1), 65-98.
- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, C.,

- Xiaocheng, F., Qin, B., & Liu, T. (2023). A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *arXiv preprint arXiv:2311.05232*.
- Hwang, W., & Salvendy, G. (2010). Number of people required for usability evaluation: The 10 ± 2 rule. *Communications of the ACM*, 53(5), 130-133.
- Infosys. (2018). *Role of AI in travel and hospitality industry*.
- Jain, A. K. (2010). Data clustering: 50 years beyond K-means. *Pattern Recognition Letters*, 31(8), 651-666.
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Chen, D., Dai, W., Chan, H. S., Madotto, A., & Fung, P. (2024). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 56(3), 1–38.
- Kandpal, N., Deng, H., Roberts, A., Wallace, E., & Raffel, C. (2023). Large language models struggle to learn long-tail knowledge. *Proceedings of the 40th International Conference on Machine Learning*, 15696–15707.
- Karpukhin, V., Oguz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., & Yih, W.-t. (2020). Dense passage retrieval for open-domain question answering. *Proceedings of EMNLP*.
- Khattab, O., & Zaharia, M. (2020). ColBERT: Efficient and accurate semantic search. *Proceedings of SIGIR*.
- Kidd, C., & Hornay, R. (2021). *Data analytics vs data analysis: What's the difference?* BMC Software. Retrieved from <https://www.bmc.com/blogs/data-analytics-vs-data-analysis>
- Kohavi, R., Tang, D., & Xu, Y. (2020). *Trustworthy online controlled experiments: A practical guide to A/B testing*. Cambridge University Press.
- Krippendorff, K. (2018). *Content analysis: An introduction to its methodology* (4th ed.). Sage publications.
- Law, R., Buhalis, D., & Cobanoglu, C. (2014). Progress on ICT in hospitality and tourism. *International Journal of Contemporary Hospitality Management*, 25(5), 727–750.
- Lemon, K. N., & Verhoef, P. C. (2016). Understanding customer experience throughout the customer journey. *Journal of Marketing*, 80(6), 69-96.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W. T., Rocktäschel, H., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474.

- Maynez, J., Narayan, S., Bohnet, B., & McDonald, R. (2020). On faithfulness and factuality in abstractive summarization. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 1906-1919.
- McHugh, M. L. (2012). Interrater reliability: the kappa statistic. *Biochemia Medica*, 276–282. <https://doi.org/10.11613/bm.2012.031>
- Micheli, M., Ponti, M., Macgilchrist, F., & Bosco, C. (2020). Emerging models of data governance in the age of datafication. *Big Data & Society*, 7(2).
- Ministry of Tourism and Sports (MOTS). (2023). *Domestic and International Tourism Statistics*. Royal Thai Government. <https://www.mots.go.th>
- Moscardo, G. (2008). Community capacity building: An emerging challenge for tourism development. In *Building community capacity for tourism development* (pp. 1-15). CABI.
- Moscardo, G. (2023). Community involvement and local knowledge in tourism development. *Journal of Sustainable Tourism*, 31(2), 345–360.
- Nielsen, J. (1994). Heuristic evaluation. In *Usability inspection methods* (pp. 25-62). John Wiley & Sons.
- OECD. (2022). *OECD Framework for the Classification of AI systems*. OECD Digital Economy Papers, No. 323, OECD Publishing, Paris.
- Ostrom, E. (1990). *Governing the commons: The evolution of institutions for collective action*. Cambridge University Press.
- Quanova Data Technologies. (2017). *Machine learning technologies for the hospitality industry*.
- Personal Data Protection Committee. (2019). *Personal Data Protection Act, B.E. 2562*. Royal Thai Government Gazette. <https://www.mdes.go.th/law/detail/3621-Personal-Data-Protection-Act-B-E--2562--2019->
- Porter, M. E., & Kramer, M. R. (2011). Creating shared value: How to reinvent capitalism—and unleash a wave of innovation and growth. *Harvard Business Review*, 89(1/2), 62-77.
- Ragas Team. (n.d.). *Ragas: Retrieval Augmented Generation Assessment framework*. Exploding Gradients. Retrieved from <https://docs.ragas.io>
- Rawte, V., Sheth, A., & Das, A. (2023). A survey of hallucination in large foundation models. *arXiv preprint arXiv:2309.05922*.
- Ricci, F., Rokach, L., & Shapira, B. (2022). *Recommender systems handbook*. Springer.
- Robinson, M. (2024). *The transformative power of generative AI in travel*. Accenture.

- Rousseeuw, P. J. (1987). Silhouettes: A graphical aid to the interpretation and validation of cluster analysis. *Journal of Computational and Applied Mathematics*, 20, 53–65.
- Salminen, J., Sengün, S., Kwak, H., Jansen, B. J., An, J., & Jung, S. G. (2018). Generating cultural personas from social data: A perspective of Middle Eastern users. *Proceedings of the 4th International Symposium on Social Networks and Social Media in the Middle East*.
- Saunders, B., Sim, J., Kingstone, T., Baker, S., Waterfield, J., Bartlam, B., Burroughs, H., & Jinks, C. (2018). Saturation in qualitative research: Exploring its conceptualization and operationalization. *Quality & Quantity*, 52(4), 1893–1907.
- Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., & Yao, S. (2023). Reflexion: Language agents with verbal reinforcement learning. *arXiv preprint arXiv:2303.11366*.
- Singhal, A. (2001). Modern information retrieval: A brief overview. *IEEE Data Engineering Bulletin*, 24(4), 35–43.
- Sivan, V. (2024). *The ultimate guide to retrieval-augmented generation*. DEV Community.
- Subramony, M., Ehrhart, M., Groth, M., Holtom, B. C., van Jaarsveld, D. D., Yagil, D., Taylor, T., Bowen, D. E., & Solnet, D. (2017). Accelerating employee-related scholarship in service management: Research streams, propositions, and a call to action. *Journal of Service Management*, 28(5), 837–865.
- Tussyadiah, I. P., et al. (2024). AI in tourism: Capabilities and limitations. *Journal of Travel Research*, 63(1), 12–28.
- UNWTO. (2023). *Tourism trends and outlook report*.
- Van Dijck, J., Poell, T., & De Waal, M. (2018). *The platform society: Public values in a connective world*. Oxford University Press.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 5998–6008.
- Wang, B., Chen, W., Pei, H., Xie, C., Kang, M., Zhang, C., Che, J., Gao, Y., Hou, L., Jian, Y., Jiao, F., Dong, Z., Shen, X., Xiong, C., Zhao, R., Zhu, C., Linden, S., Wang, H., ... Li, B. (2023). DecodingTrust: A comprehensive assessment of trustworthiness in GPT models. *Advances in Neural Information Processing Systems*, 36, 4305–4350.
- Wang, L., Ma, C., Feng, X., Zhang, Z., Yang, H., Zhang, J., Chen, Z., Tang, J., Chen, X., Lin, Y., Zhao, W. X., Wei, Z., & Wen, J.-R. (2024). A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6), 186345.

- Wei, J., He, J., Chen, K., Zhou, Y., & Tang, Z. (2017). Collaborative filtering and deep learning based recommendation system for cold start items. *Expert Systems with Applications*, 69, 29-39.
- Wen, J., Kozak, M., Yang, S., & Liu, F. (2021). COVID-19 and travel behavior. *Tourism Review*, 76(1), 74–87.
- Xi, Z., Chen, W., Guo, X., He, W., Ding, Y., Hong, B., Zhang, M., Wang, J., Jin, S., Zhou, E., Zheng, R., Fan, X., Wang, X., Xiong, L., Zhou, Y., Wang, W., Jiang, C., Zou, Y., Liu, X., ... Gui, T. (2025). The rise and potential of large language model based agents: A survey. *Science China Information Sciences*, 68(2), Article 121101.
- Xiao, L., & Wu, J. (2023). Legal document grounding with retrieval-augmented transformers. *AI and Law*, 31(2), 221–240.
- Yin, R. K. (2018). *Case study research and applications: Design and methods* (6th ed.). Sage publications.
- Yuan, H., & Hernandez, A. A. (2023). User cold start problem in recommendation systems: A systematic review. *IEEE Access*, 11, 142100–142116.
<https://doi.org/10.1109/ACCESS.2023.3341380>
- Zheng, L., Chiang, W. L., Ying, Y., Wu, S., Li, C., Zhang, Z., Lin, Z., Sun, J., Barber, J., Gallagher, R., Liang, P., Xing, E. P., & Hao, H. (2023). Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. *Advances in Neural Information Processing Systems*, 36, 46595–46623.
- Zuboff, S. (2019). *The age of surveillance capitalism: The fight for a human future at the new frontier of power*. PublicAffairs.

APPENDIX

Appendix A system data structures and knowledge corpora

A.1 Overview of the Knowledge Base

This appendix provides representative samples of the raw JSON data structures utilized to populate the NoSQL Firestore database and execute the automated evaluation phase. These structures form the foundational knowledge corpora that the contextRetriever accesses during the semantic search phase.

By structuring the data into distinct collections (User Profiles, Community Items, and Test Case Batches), the Agentic RAG system is able to programmatically cross-reference guest constraints with real-time vendor availability.

A.2 Tourist Persona Dataset (Guest Profiles)

This data structure stores the synthetic personas used during the evaluation phase, establishing the baseline demographic and temporal constraints (Length of Stay) for the Agentic Planner.

```
user_profile.json
[
  {
    "book_id": "00001",
    "guest_name": "Anong Suksamran",
    "country": "Thailand",
    "check_in_date": "2025-03-15",
    "check_out_date": "2025-03-22",
    "length_of_stay": 7
  },
  {
    "book_id": "00002",
    "guest_name": "David Miller",
    "country": "USA",
    "check_in_date": "2025-01-20",
    "check_out_date": "2025-01-25",
    "length_of_stay": 5
  },
  {
    "book_id": "00003",
    "guest_name": "Preeya Chanthavong",
    "country": "Thailand",
    "check_in_date": "2025-02-10",
    "check_out_date": "2025-02-21",
    "length_of_stay": 11
  }
]
```

A.3 Community-Based Tourism (CBT) Dataset

This data structure represents the real-time offerings of the local Koh Lanta community. It highlights the strict physical capacities (e.g., 1-4 people vs. 15-16 people) that the AI must respect to ensure operational feasibility and safety.

```
community_item.json
[
  {
    "partner": "Lanta Tourist Center",
    "item": "Four Islands by Longtail Boat (include National Park Fees)",
    "review": 5.0
  },
  {
    "partner": "Local Community",
    "item": "Private Charter Longtail Boat One Day Tour Four Islands (1-4 people)",
    "review": 4.75
  },
  {
    "partner": "Lanta Thai Cookery School",
    "item": "Dinner Course",
    "review": 5.0
  },
  {
    "partner": "Koh Lanta Community Museum",
    "item": "Donation for the Koh Lanta Community Museum",
    "review": 5.0
  }
]
```

A.4 Automated Genkit Evaluation Inputs (Test Case Batches)

This data structure defines the programmatic inputs passed to the Firebase Genkit evaluation flow. The dataset was processed in batches (e.g., A-000001.json). Each record acts as a simulated prompt, linking a specific guest profile (book_id) to a desired set of activities, thereby triggering the Agentic Planner to retrieve constraints and generate a localized itinerary.

```
A-000001.json
[
  {
    "testCaseId": "e41fe0a0-ba7b-4e19-aa13-2cc1575b9a8e",
    "input": {
      "hotel_id": "000004",
      "book_id": "00001",
      "activity": "Explore the Islands by Longtail Boat, Explore Local Markets"
    }
  },
  {
    "testCaseId": "99bed8a2-2449-4e75-a27a-d953084e62bc",
    "input": {
      "hotel_id": "000004",
      "book_id": "00002",
      "activity": "Visit a Museum, Take a Cooking Class"
    }
  }
]
```



```
},  
{  
  "testCaseId": "57d69cd4-38a3-4f3c-8621-f366c01d2039",  
  "input": {  
    "hotel_id": "000004",  
    "book_id": "00003",  
    "activity": "Visit a Museum,Use local transfer services"  
  }  
}  
]
```

Appendix B: Ragas Framework & Human Expert Evaluation Metrics

B.1 Overview of the Ragas Scoring Dataset

Overview of Evaluation Data

This appendix details the raw evaluation metrics of the Agentic RAG (Firebase Genkit) architecture. Testing was conducted using 150 simulated queries derived from 50 synthetic traveler personas (detailed in Appendix A).

To ensure rigorous validation, every system output was subjected to a dual-evaluation methodology. Outputs were scored simultaneously by the automated Ragas framework and validated by local human domain experts. Both evaluators scored the outputs across three critical dimensions:

Maliciousness: The presence of harmful, unsafe, or physically dangerous itinerary recommendations (Lower is better; 0 indicates perfect safety).

Answer Relevancy: The semantic alignment between the traveler's specific constraints (e.g., time of day, group size) and the generated response (Higher is better).

Faithfulness: The factual accuracy of the response, measuring whether the AI successfully grounded its recommendation exclusively in the live community database without hallucinating (Higher is better; 1 indicates zero hallucination).

The data below demonstrates the correlation between the automated AI-as-a-judge metrics (Ragas) and the ground-truth scoring provided by human domain experts.

B.2 Dual-Evaluation Query Metrics (Sample)

	A	B	C	D	E	F	G	H	I	J	K	
1						Genkit			Base LLM			
2	No	Hotel ID	Booking ID	Activity	Maliciousness	Answer Relevan	Faithfulness	Maliciousness	Answer Relevan	Faithfulness		
3	1	000004	00001	Explore the Islands by Longtail Boat,Explic	0.00	1.00	0.81	0.00	0.80	0.86		
4	2	000004	00002	Visit a Museum,Take a Cooking Class	0.00	0.70	0.90	0.00	0.50	0.25		
5	3	000004	00003	Visit a Museum,Use local transfer services	0.00	0.77	0.50	0.00	1.00	1.00		
6	4	000004	00004	Explore the Islands by Longtail Boat	0.00	0.76	1.00	0.00	0.50	0.50		
7	5	000004	00005	Explore Local Markets,Visit a Museum	0.00	1.00	1.00	0.00	0.50	0.25		
8	6	000004	00006	Take a Cooking Class	0.00	0.00	0.30	0.00	1.00	1.00		
9	7	000004	00007	Use local transfer services,Visit a Museum	0.00	0.73	0.90	0.00	0.50	0.25		
10	8	000004	00008	Explore the Islands by Longtail Boat	0.00	0.33	0.93	0.00	0.50	0.25		
11	9	000004	00009	Take a Cooking Class,Explore Local Mark	0.00	0.75	0.70	0.00	0.50	0.25		
12	10	000004	00010	Visit a Museum	0.00	0.28	1.00	0.00	0.50	0.25		
13	11	000004	00011	Explore the Islands by Longtail Boat,Use l	0.00	0.93	0.50	0.00	0.50	0.25		
14	12	000004	00012	Explore Local Markets	0.00	0.92	0.50	0.00	0.50	0.25		
15	13	000004	00013	Take a Cooking Class,Visit a Museum	0.00	0.27	0.90	0.00	0.50	0.25		
16	14	000004	00014	Use local transfer services	0.00	0.66	0.60	0.00	0.50	0.25		
17	15	000004	00015	Explore the Islands by Longtail Boat,Explic	0.00	0.82	0.24	0.00	0.50	0.25		
18	16	000004	00016	Take a Cooking Class	0.00	0.10	0.10	0.00	0.50	0.25		
19	17	000004	00017	Visit a Museum,Use local transfer services	0.00	0.29	0.50	0.00	0.50	0.25		

	A	B	C	D	E	F	G	H	I	J	K	
1						Genkit			Human			
2	No	Hotel ID	Book ID	Request	Maliciousness	Answer Relevan	Faithfulness	Maliciousness	Answer Relevan	Faithfulness	Result	
3	1	000004	00001	I want to experience a gondola ride. We p	0.00	0.64	1.00	0.00	0.59	0.50	"For a magical mornin	
4	2	000004	00002	I'm looking for a free or low-cost cultural a	0.00	0.87	0.40	0.00	0.84	0.75	"For a relaxed and cul	
5	3	000004	00003	We want to visit Koh Rok for a day trip. W	0.00	0.88	0.67	0.00	0.91	0.67	"From Koh Lanta, you	
6	4	000004	00004	Can we book a gondola trip? We want to g	0.00	0.78	1.00	0.00	0.61	1.00	"Yes, for an enchantin	
7	5	000004	00005	We are looking for a private boat charter. I	0.00	0.84	0.60	0.00	0.81	0.50	"For your group of 11 i	
8	6	000004	00006	Is there a boat trip to Koh Ha for snorkelin	0.00	0.79	1.00	0.00	0.82	1.00	"Yes, in Koh Lanta, joi	
9	7	000004	00007	We'd like to take a relaxing boat ride in the	0.00	0.72	1.00	0.00	0.00	1.00	"For a relaxing mornin	
10	8	000004	00008	Do you have a private longtail boat availat	0.00	0.83	1.00	0.00	0.54	0.83	"Yes, a private longtail	
11	9	000004	00009	I'm looking for a tour with crystal clear wat	0.00	0.74	0.00	0.00	0.80	1.00	"Koh Rok is perfect for	
12	10	000004	00010	We want to avoid the crowds and go priva	0.00	0.91	0.00	0.00	0.77	0.60	"To avoid crowds in Kc	
13	11	000004	00011	Is there a place where I can learn about th	0.00	0.92	0.75	0.00	0.77	1.00	"Yes, visit the Koh Lan	
14	12	000004	00012	My partner and I are looking for a romanti	0.00	0.82	0.50	0.00	0.75	1.00	"For a romantic dinner	
15	13	000004	00013	We are a group of friends wanting to visit i	0.00	0.81	0.89	0.00	0.79	0.67	"Yes, join the 4 Island	
16	14	000004	00014	Can we book a private tour? We are a gro	0.00	0.78	0.50	0.00	0.74	0.75	"Yes, book a private lo	
17	15	000004	00015	I want to experience a traditional gondola	0.00	0.77	0.00	0.00	0.90	0.00	"Experience a traditior	
18	16	000004	00016	We prefer a private tour over a shared one	0.00	0.00	0.91	0.00	0.80	0.75	"Private longtail boat c	
19	17	000004	00017	Do you have any special dinner courses fo	0.00	0.74	0.33	0.00	0.84	0.75	"For celebrations in Kc	

Appendix C: System Prompt Configurations (Dotprompts)

C.1 Overview of Agentic Instructions

This appendix details the exact programmatic instructions (prompts) utilized within the Firebase Genkit environment to orchestrate the Agentic RAG system. By defining strict system directives, input schemas, and output formats within .prompt files, the architecture restricts the Large Language Model's parametric memory and forces adherence to local community constraints.

C.2 Generation Phase: The Recommendation Prompt

This Dotprompt is executed after the Context Retriever has successfully fetched the relevant community data. It instructs the Gemini model to synthesize the final itinerary while explicitly forbidding hallucinations and enforcing safety guardrails.

```
You are a helpful AI product recommender. Your task is to select a few items from the
list of these partners:
{{#each partners}}
  partner id: {{partner_id}}, partner name: {{partner_name}}
{{/each}}
And consider the following partner's items:
{{#each items}}
  item id: {{item_id}}, item name: {{item_name}}, item description: {{item_description}}
{{/each}}
Consider the user's past order history:
{{#each orders}}
  order id: {{order_id}}, partner id: {{partner_id}}, item id: {{item_id}}, item name:
{{item_name}}
{{/each}}
And the following item review scores (on a scale of 1 to 5):
{{#each reviews}}
  item id: {{item_id}}, score: {{review}}
{{/each}}
And consider the user's past viewing history:
{{#each views}}
  item id: {{item_id}}
{{/each}}
And consider the user's activity:
{{#each activity}}
```

```

    {{activity_name}}
  {{/each}}
  Based on this information, recommend a few products that the user might be interested in.
  **Optional Enhancements to Guide Recommendations:**
  * **Number of Recommendations:** Please recommend about 5 items.
  * **Prioritization based on Orders:** Focus on recommending items that the user hasn't
  ordered before but are related to their past purchases (either the same item or similar
  items).
  * **Prioritization based on Reviews:** Prioritize recommending items with a review score
  above a certain threshold (e.g., recommend items with a 'review' score of 4 or higher).
  * **Reasoning:** Briefly explain why you are recommending each item, referencing either
  their past orders or positive review scores.
  * **Consider Item Description:** The 'item_descption' can provide more context for making
  relevant recommendations. Consider similarities in descriptions.
  \

```

C.3 Retrieval Phase: The Semantic Router Prompt

This Dotprompt serves as the Agentic decision-maker before the database search occurs. It analyzes the user's raw query and determines the optimal metadata tags needed to search the NoSQL Vector Database efficiently.

```

const productRecommendFlow = ai.defineFlow({
  name: "productRecommendFlow",
  inputSchema: z.object({
    hotel_id: z.string(),
    book_id: z.string(),
    activity: z.string().optional(),
  }),
},
async (input) => {
  let activityData: any[] = [];
  if (input.activity) {
    activityData = input.activity.split(',').map(s => ({ activity_name: s.trim() }));
  }
  const partners = await firestore.collection(`hotels/${input.hotel_id} +
  `/partner`).get();
  const partnerDataPromises = partners.docs.map(async (doc) => {
    const data = doc.data() as any;
    return {
      partner_id: doc.id,
      partner_name: data.partner_name,
      ...data,
    };
  });
  const partnerData = await Promise.all(partnerDataPromises);

```

```

const queryText = input.activity || "recommended items for hotel guests";
const userEmbedding = await ai.embed({
  embedder: geminiEmbedding001,
  content: queryText,
});

let allPartnerItems: any[] = [];

for (const partner of partnerData) {
  const partnerItemsColl =
    firestore.collection(`partners/${partner.partner_id}/items`);
  const vectorQuery = partnerItemsColl.findNearest({
    vectorField: 'embedding',
    queryVector: userEmbedding[0].embedding,
    distanceMeasure: 'COSINE',
    limit: 5,
  });
  const snapshot = await vectorQuery.get();
  const items = snapshot.docs.map(doc => ({ item_id: doc.id, ...doc.data() }));
  allPartnerItems = [...allPartnerItems, ...items];
}

const orders = await firestore.collectionGroup(`orders`).where(`hotel_id`, `==`,
input.hotel_id).where(`book_id`, `==`, input.book_id).get();
const result2 = orders.docs.map((doc) => {
  const data = doc.data() as any;
  return {
    order_id: doc.id,
    ...data,
  }
});

const reviews = await firestore.collectionGroup(`reviews`).get();
const result3 = reviews.docs.map((doc) => doc.data() as any).filter((data) =>
  allPartnerItems.some((item) => item.item_id === data.item_id)
).map((data) => ({
  ...data,
}));

const views = await firestore.collectionGroup(`views`).get();
const result4 = views.docs.reduce((acc: any[], doc) => {
  const data = doc.data() as any;
  if (data.book_id === input.book_id) {
    acc.push({
      ...data,
    });
  }
  return acc;
}, []);

const factDocs = await ai.retrieve({

```

```
    retriever: contextRetriever,  
    query: '',  
  });  
  
  const { output } = await recommendDotprompt({  
    items: allPartnerItems,  
    partners: partnerData,  
    orders: result2,  
    reviews: result3,  
    views: result4,  
    activity: activityData  
  }, { context: factDocs });  
  
  return output;  
}  
);  
  
export const productRecommend = onCallGenkit({  
  secrets: isLocal ? [] : [apiKey],  
  cors: true,  
}, productRecommendFlow);
```

Appendix D: Nosql Database Schema And Data Dictionary

D.1 Overview of the Firestore Architecture

This appendix details the schema design for the NoSQL database utilized in the Tripartite Application Ecosystem. While NoSQL environments (such as Firebase Firestore) allow for flexible document structures, a strict schema was enforced at the application level to ensure data integrity during the Retrieval-Augmented Generation (RAG) process.

The database is divided into three primary collections: hotels, guests and partners. Notably, the schema includes high-dimensional vector fields required for semantic similarity search.

D.2 Collection: hotels

```
export const HotelAttributesMapping = {
  hotel_name: 'Hotel Name',
  location: 'GPS Location',
  hotel_info: 'Hotel Info',
  hotel_image: 'Hotel Image',
  sla: 'Hotel Service Time',
};

export interface Hotel {
  hotel_name: string;
  location: string;
  hotel_info: string;
  hotel_image: string;
  sla: number;
}

export const RoomtypeAttributesMapping = {
  roomtype_name: 'RoomType Name',
  roomtype_description: 'RoomType Description',
  image_url: 'RoomType Image',
};

export interface Roomtype {
  roomtype_name: string;
  roomtype_description: string;
  image_url: string;
}

export const RoomAttributesMapping = {
```



```
roomtype_name: 'RoomType',
room_name: 'Room Name',
room_description: 'Room Description',
room_price: 'Room Price',
image_url: 'Room Image',
};
```

D.3 Collection: guests

```
export class Guest {
  id?: string;
  book_number?: string;
  first_name?: string;
  last_name?: string;
  email?: string;
  country_code?: string;
  phone_no?: string;
  id_image?: string;
  signature_image?: string;
  check_in?: Date;
  check_out?: Date;
  created_at?: Date;
  updated_at?: Date;
}
```

D.4 Collection: partners

```
export class HotelPartner {
  id?: string;
  partner_id?: string;
  partner_name?: string;
  partner_image?: string;
  partner_distance?: number;
  partner_status?: string;
  created_at?: Date;
  updated_at?: Date;
}
```

```
export class Service {
  id?: string;
  servicetype_id?: string;
  servicetype_ref?: DocumentReference;
  service_name?: string;
  service_description?: string;
```

```
service_sla?: number;  
service_price?: number;  
image_name?: string;  
image_url?: string;  
status?: string;  
created_at?: Date;  
updated_at?: Date;  
}
```

CURRICULUM VITAE

NAME	Mr. Somatat Na Takuatung
DATE OF BIRTH	15 April 1966
EDUCATIONAL RECORD	
BACHELOR'S DEGREE	Bachelor of Science (Applied Statistics) King Mongkut's University of Technology North Bangkok
MASTER'S DEGREE	Master of Design (Human-Centered Design) Institute of Design Illinois Institute of Technology, 2000.
SCHOLARSHIP	Lab-Based, School of Architecture and Design
EMPLOYMENT RECORD	Cyber Works Media Group, Chicago, 2000 - 2004 Soft Square Group, 2004 - 2008 Soft2, 2008 - Present
PUBLICATION	na takuatung, somatat, 2023. "Boutique Hotel Service Digitalization: A Business Owner Study." Journal of Architectural/Planning Research and Studies (JARS) 21 (1): 167-184. https://doi.org/10.56261/jars.v21.256845 .

King Mongkut's University of Technology Thonburi
Agreement on Intellectual Property Rights Transfer for Postgraduate Students

Date.....2 August 2026.....

Name.....**Mr. Somatat**.....
Surname/Family Name.....**Na Takuatung**.....
Student Number...**63120800003**... who is a student of King Mongkut's University of
Technology Thonburi (KMUTT) in ☐ Graduate Diploma ☐ master's degree
☐ Doctoral Degree
Program**Doctor of Philosophy**..... Field of Study..... **Design and Planning**...
Faculty/School.....**School of Architecture and Design**.....
Home Address**254/2 Krungthep-Nonthaburi Rd Bang Sue Dusit Bangkok** ...
Postal Code.....**10800**.....Country.....**THAILAND**.....

I, as ‘Transferer’, hereby transfer the ownership of my thesis copyright to King Mongkut’s University of Technology Thonburi who has appointed **Asst. Prof. Waraluk Pansuwan, Dean of School of Architecture and Design** to be ‘Transferee’ of copyright ownership under the ‘Agreement’ as follows.

1. I am the author of the thesis entitled **The agentic ai revolution in sustainable tourism: a Retrieval-Augmented Generation (RAG) framework for community-based knowledge hubs in Thailand** under the supervision of **Assoc. Prof. Chokeanand Bussracumpakorn, Ph. D.** who is my supervisor, in accordance with the Thai Copyright Act B.E. 2537. The thesis is a part of the curriculum of KMUTT.

2. I hereby transfer the copyright ownership of all my works in the thesis to KMUTT throughout the copyright protection period in accordance with the Thai Copyright Act B.E. 2537, effective on the approval date of thesis proposal consented by KMUTT.

3. To have the thesis distributed in any form of media, I shall in each and every case stipulate the thesis as the work of KMUTT.

4. For my own distribution of thesis or the reproduction, adjustment, or distribution of thesis by the third party in accordance with the Thai Copyright Act B.E. 2537 with remuneration in return, I am subject to obtain a prior written permission from KMUTT.

5. To use any information from my thesis to make an invention or create any intellectual property works within ten (10) years from the date of signing this Agreement, I am subject to obtain prior written permission from KMUTT, and KMUTT is entitled to have intellectual property rights on such inventions or intellectual property works, including entitling to take royalty from licensing together with the distribution of any benefit deriving partly or wholly from the works in the future, conforming with the Regulation of King Mongkut's Institute of Technology Thonburi *Re* the Administration of Benefits deriving from Intellectual Property B.E. 2538.

6. If the benefits arise from my thesis or my intellectual property works owned by KMUTT, I shall be entitled to gain the benefits according to the allocation rate stated

in the Regulation of King Mongkut's Institute of Technology Thonburi *Re* the
Administration of Benefits deriving from Intellectual Property B.E. 2538.

Signature.....Transferor
(Mr. Somatat Na Takuatung)
Student

Signature.....Transferee
(Asst. Prof. Waraluk Pansuwan)
Dean of School of Architecture and Design

Signature.....Witness
(Assoc. Prof. Dr. Chokeanand Bussracumpakorn)

Signature.....Witness
(Prof. Dr. Pattanasak Mongkolwat)